

Can Machine Learning Variable Selection and Individual Commodity Fundamentals Improve Futures Return Forecasting?

Leo Specks

Supervisor: Prof. Massimo Guidolin
MSc Finance, Università Bocconi

July 2027

Last Modified: Monday 15th June, 2026 (01:29)

Download latest version [here](#).

Contents

1	Test Chapter	3
1.1	(Sub)plots and Tables	3
1.2	Citations	6
1.3	Codeblock	7
2	Introduction	8
2.1	Motivation & Background	8
2.2	The Research Gap	8
2.3	Research Questions & Hypotheses	8
2.4	Contributions to the Literature	8
2.5	Thesis Outline	8
3	Literature Review	9
3.1	Theories of Commodity Futures Returns	9
3.2	Prevailing Commodity Forecasting Approaches	9
3.3	Cross-Market Transferability	9
3.3.1	Domain Adaptation and Fine-Tuning	9
4	Methodology	10
4.1	Baseline Framework (Replication)	10
4.2	Extension I: Penalized Regression	11
4.3	Extension II: Individual Fundamental Predictors	11
4.4	OOS Evaluation Criteria	11
4.5	Portfolio Back-Testing Framework	11
5	Data	12
5.1	Commodity Futures Returns	12
5.2	Macroeconomic Factors	12
5.3	Aggregate Commodity-Specific Factors	12
5.4	Individual Fundamental Predictors	12
5.5	Sample Splits and Descriptive Statistics	12
6	Replication & Baseline Results	13
6.1	In-Sample Estimation (1989-2003)	13
6.2	OOS Predictive Accuracy (2004-2012)	13
6.3	Portfolio Back-Testing (2004-2012)	13
6.4	Replication Assessment	13
7	Extension I: ML Variable Selection	14
7.1	Model Specification and Implementation	14
7.2	Variable Selection Patterns	14
7.3	OOS Forecasting Results - Original Sample (2004-2012)	14
7.4	OOS Forecasting Results - Extended Sample (2013-2024)	14
7.5	Portfolio Back-Testing	14
7.6	Robustness Checks	14

8	Extension II: Individual Commodity Fundamentals	15
8.1	Predictor Construction by Commodity Group	15
8.2	Individual vs. Aggregate Factors: Are They Substitutes?	15
8.3	OOS Forecasting Results (2004-2024)	15
8.4	Combined Model	15
8.5	Portfolio Back-Testing	15
8.6	Regime Interaction	15
9	Synthesis & Discussion	16
9.1	Model Horse Race	16
9.2	Statistical vs. Economic Loss Functions Revisited	16
9.3	Implications for Commodity Market Segmentation	16
9.4	Practical Implications	16
10	Conclusions	17
10.1	Summary of Findings	17
10.2	Limitations	17
10.3	Future Research Directions	17
	References	19
	List of Acronyms	20

1 Test Chapter


Etiam gravida varius ex quis semper. Morbi molestie ipsum id purus faucibus, laoreet cursus tortor sollicitudin. In sodales blandit nunc, vitae ultricies orci bibendum non. Sed dapibus ullamcorper magna, vitae tempor dolor condimentum id. In posuere ex id vehicula tempus. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos himenaeos. Aenean auctor justo eget tincidunt condimentum. Integer euismod massa ligula, at imperdiet nisi lobortis in.

1.1 (Sub)plots and Tables

Nunc volutpat, massa eu dignissim dignissim, purus nunc vestibulum libero, quis bibendum nulla diam sed risus. Mauris elementum turpis nec massa mattis dictum. Nulla ac dolor eu neque consequat suscipit. Fusce sed lorem tellus. Integer efficitur orci sed ligula vestibulum, at sodales odio eleifend. Suspendisse luctus diam at nisi pharetra sollicitudin. Proin est sapien, efficitur vitae sodales sodales, pellentesque in tortor. Duis at commodo diam. Morbi lobortis dictum justo vitae suscipit. Vivamus dictum lacus velit, vitae sagittis lorem sodales eu. Nam a arcu rutrum, aliquam mauris quis, tempus ante. Sed vehicula tortor id arcu dapibus, quis blandit dui egestas. Mauris in ullamcorper augue. Aenean vehicula neque quis felis lacinia, ut pharetra felis vestibulum. Maecenas sit amet lacus scelerisque orci condimentum pellentesque et vehicula lacus. Duis vitae ex a risus molestie pretium. In hac habitasse platea dictumst. Phasellus fringilla sit amet erat quis euismod. Vivamus turpis eros, vulputate id condimentum sit amet, tempus a velit. Interdum et malesuada fames ac ante ipsum primis in faucibus. Donec et convallis tortor. Aenean porttitor fermentum urna ut condimentum.

Figure 1

Example figure with two subfigures.




Integer luctus in Panel 1a, ante eu finibus facilisis in Figure, turpis dui scelerisque sapien, quis elementum dui tortor vitae nisl. Aliquam mattis mauris in dui semper lacinia.

Cras eros tellus, condimentum sed hendrerit quis, vehicula quis risus. In commodo in orci at sodales. Etiam gravida varius ex quis semper. The entire model used 50,000 samples. The method was 253.4 times faster. This corresponds to 37.1% of the data. Morbi molestie ipsum id purus faucibus, laoreet cursus tortor sollicitudin. In sodales blandit nunc, vitae ultricies orci bibendum non. Sed dapibus ullamcorper magna, vitae tempor dolor condimentum id. In posuere ex id vehicula tempus. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos himenaeos. Aenean auctor justo eget tincidunt condimentum. Integer euismod massa ligula, at imperdiet nisi lobortis in. Nunc porttitor vehicula diam iaculis mattis. Interdum et malesuada fames ac ante ipsum primis in faucibus. Pellentesque posuere ligula id felis vehicula tincidunt [1]. Mauris et mi ex. Morbi nec ligula mollis, consectetur felis id, auctor leo. Vivamus a molestie ante. Suspendisse placerat dolor ex. Nam interdum luctus felis eu eleifend. Mauris ac vulputate ex, id feugiat ipsum. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut ex leo, ullamcorper ut magna eget, tempus consectetur orci. Nunc vitae ante lectus. Ut malesuada diam ut arcu condimentum, at maximus justo ullamcorper. Mauris vehicula tortor in dolor feugiat elementum. Pellentesque tellus ligula, condimentum et felis et, ultrices laoreet turpis. Aliquam quis dolor luctus, iaculis diam vel, sodales dui [2].

Figure 2

Example Figure (width should be 16cm)



Mauris ultricies arcu ut neque semper, et pretium arcu sagittis. Mauris vestibulum maximus arcu. Sed vulputate convallis bibendum. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Suspendisse posuere arcu sed ultricies gravida. In hac habitasse platea dictumst. Aliquam id erat id nunc luctus bibendum. Curabitur tincidunt leo nec urna posuere vestibulum. Morbi lacinia aliquam arcu, quis interdum tortor congue id. Aliquam semper varius imperdiet.

Donec tincidunt magna et dolor volutpat aliquet. Etiam ac pretium ante. Nam quis lacus sit amet metus egestas laoreet. Sed in orci venenatis, laoreet mi non, ullamcorper quam. Nullam eu lobortis nibh. Duis id neque vitae dolor blandit maximus. Vivamus ligula velit, posuere ac tortor non, euismod aliquam nisi. Etiam a ultrices nisi, vitae tincidunt

leo. Donec scelerisque felis magna, sit amet euismod ex sodales a. Nunc consectetur purus nunc, vitae volutpat lectus fringilla vitae. Suspendisse ut pharetra justo, non ultrices est. Fusce sit amet ante non sapien mollis elementum id eu augue. Donec consectetur risus eget tellus pellentesque, at maximus mi elementum. Morbi id dapibus ligula. Praesent vel tempor ipsum. Phasellus at justo venenatis, lobortis leo vitae, molestie metus. Cras ullamcorper ante eu mi laoreet, quis varius tellus consectetur¹.

Donec ultrices consectetur dolor sed sodales. Cras efficitur turpis vitae lacus mollis cursus. Pellentesque iaculis nulla ac vestibulum fermentum. Pellentesque et velit non turpis suscipit congue. Aliquam vitae purus a nibh cursus sagittis. Nullam egestas viverra purus vitae varius. Sed maximus tempor vestibulum. Integer luctus in Panel 1a, ante eu finibus facilisis in Figure, turpis dui scelerisque sapien, quis elementum dui tortor vitae nisl. Aliquam mattis mauris in dui semper lacinia. Cras eros tellus, condimentum sed hendrerit quis, vehicula quis risus. In commodo in orci at sodales. Etiam gravida varius ex quis semper. The entire model used 50,000 samples. The method was 253.4 times faster. This corresponds to 37.1% of the data. Morbi molestie ipsum id purus faucibus, laoreet cursus tortor sollicitudin. In sodales blandit nunc, vitae ultricies orci bibendum non. Sed dapibus ullamcorper magna, vitae tempor dolor condimentum id. In posuere ex id vehicula tempus. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos himenaeos. Aenean auctor justo eget tincidunt condimentum. Integer euismod massa ligula, at imperdiet nisi lobortis in. Nunc porttitor vehicula diam iaculis mattis. Interdum et malesuada fames ac ante ipsum primis in faucibus. Pellentesque posuere ligula id felis vehicula tincidunt [1]. Mauris et mi ex. Morbi nec ligula mollis, consectetur felis id, auctor leo. Vivamus a molestie ante. Suspendisse placerat dolor ex. Nam interdum luctus felis eu eleifend. Mauris ac vulputate ex, id feugiat ipsum. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut ex leo, ullamcorper ut magna eget, tempus consectetur orci. Nunc vitae ante lectus. Ut malesuada diam ut arcu condimentum, at maximus justo ullamcorper. Mauris vehicula tortor in dolor feugiat elementum. Pellentesque tellus ligula, condimentum et felis et, ultrices laoreet turpis. Aliquam quis dolor luctus, iaculis diam vel, sodales dui [2]. Integer luctus in Panel 1a, ante eu finibus facilisis in Figure, turpis dui scelerisque sapien, quis elementum dui tortor vitae nisl. Aliquam mattis mauris in dui semper lacinia. Cras eros tellus, condimentum sed hendrerit quis, vehicula quis risus. In commodo in orci at sodales. Etiam gravida varius ex quis semper. The entire model used 50,000 samples. The method was 253.4 times faster. This corresponds to 37.1% of the data. Morbi molestie ipsum id purus faucibus, laoreet cursus tortor sollicitudin. In sodales blandit nunc, vitae ultricies orci bibendum non. Sed dapibus ullamcorper magna, vitae tempor dolor condimentum id. In posuere ex id vehicula tempus. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos himenaeos. Aenean

¹This would be the first footnote.

auctor justo eget tincidunt condimentum. Integer euismod massa ligula, at imperdiet nisi lobortis in. Nunc porttitor vehicula diam iaculis mattis. Interdum et malesuada fames ac ante ipsum primis in faucibus. Pellentesque posuere ligula id felis vehicula tincidunt [1]. Mauris et mi ex. Morbi nec ligula mollis, consectetur felis id, auctor leo. Vivamus a molestie ante. Suspendisse placerat dolor ex. Nam interdum luctus felis eu eleifend. Mauris ac vulputate ex, id feugiat ipsum. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut ex leo, ullamcorper ut magna eget, tempus consectetur orci. Nunc vitae ante lectus. Ut malesuada diam ut arcu condimentum, at maximus justo ullamcorper. Mauris vehicula tortor in dolor feugiat elementum. Pellentesque tellus ligula, condimentum et felis et, ultrices laoreet turpis. Aliquam quis dolor luctus, iaculis diam vel, sodales dui [2]. Integer luctus in Panel 1a, ante eu finibus facilisis in Figure, turpis dui scelerisque sapien, quis elementum dui tortor vitae nisl. Aliquam mattis mauris in dui semper lacinia. Cras eros tellus, condimentum sed hendrerit quis, vehicula quis risus. In commodo in orci at sodales. Etiam gravida varius ex quis semper. The entire model used 50,000 samples. The method was 253.4 times faster. This corresponds to 37.1% of the data. Morbi molestie ipsum id purus faucibus, laoreet cursus tortor sollicitudin. In sodales blandit nunc, vitae ultricies orci bibendum non. Sed dapibus ullamcorper magna, vitae tempor dolor condimentum id. In posuere ex id vehicula tempus. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos himenaeos. Aenean auctor justo eget tincidunt condimentum. Integer euismod massa ligula, at imperdiet nisi lobortis in. Nunc porttitor vehicula diam iaculis mattis. Interdum et malesuada fames ac ante ipsum primis in faucibus. Pellentesque posuere ligula id felis vehicula tincidunt [1]. Mauris et mi ex. Morbi nec ligula mollis, consectetur felis id, auctor leo. Vivamus a molestie ante. Suspendisse placerat dolor ex. Nam interdum luctus felis eu eleifend. Mauris ac vulputate ex, id feugiat ipsum. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut ex leo, ullamcorper ut magna eget, tempus consectetur orci. Nunc vitae ante lectus. Ut malesuada diam ut arcu condimentum, at maximus justo ullamcorper. Mauris vehicula tortor in dolor feugiat elementum. Pellentesque tellus ligula, condimentum et felis et, ultrices laoreet turpis. Aliquam quis dolor luctus, iaculis diam vel, sodales dui [2]. Nunc tempor nulla justo, a posuere lorem tincidunt eget. Praesent dictum sed eros in pharetra. Vivamus nec dictum risus. Integer ac massa at magna rutrum pulvinar eget nec nunc. Curabitur tempor at lorem vitae venenatis. Suspendisse auctor posuere iaculis. Suspendisse egestas lobortis diam. Integer mattis. The approach by Guidolin and Pedio [2] has shown that this is realistic.

1.2 Citations

This is a \cite{}: [2], this a \textcite{}: Guidolin and Pedio [2], this a \parencite{}: [2], this a \citeauthor{}: Guidolin and Pedio, this a \citeyear{}: 2021, and this a

parencite with two references: [1, 2].

Table 1
Citation Styles with BibLaTeX

Command	authoryear	numeric
<code>\cite{}</code>	Author, Year	[index]
<code>\textcite{}</code>	Author (Year)	Author [index]
<code>\parencite{}</code>	(Author, Year)	[index]
<code>\citeauthor{}</code>	Author	Author
<code>\citeyear{}</code>	Year	Year (not linked)

Use `\parencite{}` at the end of a statement and `\textcite{}` when referencing in the middle of a sentence.

1.3 Codeblock

The algorithm can be viewed in Codeblock 1. Nullam egestas viverra purus vitae varius. Sed maximus tempor vestibulum. Nunc tempor nulla justo, a posuere lorem tincidunt eget. Praesent dictum sed eros in pharetra. Vivamus nec dictum risus. Integer ac massa at magna rutrum pulvinar eget nec nunc. Curabitur tempor at lorem vitae venenatis. Suspendisse auctor posuere iaculis. Suspendisse egestas lobortis diam. Integer mattis.

Code 1

Example Codeblock

```
# KPI OVERVIEW
kpis = pd.DataFrame({
    "Ann. Return (approx)": ann_return_simple,
    "Ann. Return (compounding)": ann_return,
    "Ann. Volatility": ann_vol,
    "Sharpe": sharpe,
    "CAGR": cagr,
    "Max Drawdown": max_dd
})

kpis.T
```


2 Introduction

2.1 Motivation & Background

2.2 The Research Gap

2.3 Research Questions & Hypotheses

2.4 Contributions to the Literature

2.5 Thesis Outline

3 Literature Review

3.1 Theories of Commodity Futures Returns

3.2 Prevailing Commodity Forecasting Approaches

3.3 Cross-Market Transferability

Some of the more complex models seen in Subsection 3.2 suffer from needing of a lot of data. For some more niche markets with shorter histories and more constrained datasets this raises a problem. The most straightforward and perhaps naive solution is to simply train models on data-rich markets and implement them on data-sparse ones. However, it has been shown that traditional models fail when being transferred cross-market. When trying to transfer a convolutional neural network (CNN) model trained on US commodity futures to predict return of the Chinese market, Hao et al. [3] find that this model fails to systematically outperform models retrained on data specific to the Chinese market. These results suggest that there are no significant return-predictive patterns from US commodity futures which generalize across different markets. Hao et al. [3] give two explanations on why employing models cross-market fails in commodity markets. **insert reasons** Furthermore, Nguyen and Walther [4] have also found results to be heterogeneous across different commodity futures as no single model fit all commodities. It seems that traditional models lack cross-market transferability across both different countries and commodities.

This shortcoming naturally invites the exploration of methodologies, which aim to built predictors, that allow for differences cross-market. The goal is to leverage data-rich markets, while still being able to include market-specific characteristics. This approach can be categorised as a transfer learning technique, which is a subset of machine learning, that allows for the assumption of different domains, tasks and distributions between the training and test set [5].

3.3.1 Domain Adaptation and Fine-Tuning

Mashetty et al. [6] This approach proved to outperform both a traditional machine learning model and a deep learning model, which were trained only on target data, in all three chosen metrics, namely mean absolute error (MAE), root mean squared error (RMSE) and directional accuracy.

4 Methodology

4.1 Baseline Framework (Replication)

We begin by gathering all T observations of n variables into a data matrix $\mathbf{X} \in \mathbb{R}^{T \times n}$, where each row is a month and each column is a macroeconomic variable. Each column is demeaned and standardized to unit variance using the StandardScaler:

$$X_{\text{scaled},ij} = \frac{X_{\text{raw},ij} - \mu_j}{\sigma_j}. \quad (1)$$

Let $\Sigma = V[\mathbf{X}] \in \mathbb{R}^{n \times n}$ denote the variance-covariance matrix of $\mathbf{X}$, which is calculated as:

$$\Sigma = \frac{1}{T-1} \mathbf{X}'\mathbf{X}, \quad (2)$$

and let $\Sigma = \mathbf{U}\mathbf{\Lambda}\mathbf{U}'$ be its spectral decomposition, where $\mathbf{U} \in \mathbb{R}^{n \times n}$ is an orthogonal matrix of order n , whose columns are the eigenvectors $\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_n$ of Σ , and $\mathbf{\Lambda} \in \mathbb{R}^{n \times n}$ is a diagonal matrix of the same order, whose diagonal entries are the corresponding eigenvalues $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_n$, which, by convention, are sorted from the largest to the smallest. By multiplying the original scaled data matrix $\mathbf{X}$ with $\mathbf{U}$, we derive the principal components of $\mathbf{X}$ as the elements of the matrix $\mathbf{F} \in \mathbb{R}^{T \times n}$, or formally:

$$\mathbf{F} = \mathbf{X}\mathbf{U}. \quad (3)$$

To assess how much of the total variance in the original n variables is captured by the k -th principal component, or alternatively by the first K principal components, we can compute the variance-covariance matrix of $\mathbf{F}$ as:

$$\begin{aligned} V[\mathbf{F}] &= V[\mathbf{X}\mathbf{U}] \\ &= \mathbf{U}'\Sigma\mathbf{U} \\ &= \mathbf{U}'\mathbf{U}\mathbf{\Lambda}\mathbf{U}'\mathbf{U} \\ &= \mathbf{I}_n\mathbf{\Lambda}\mathbf{I}_n \\ &= \mathbf{\Lambda}, \end{aligned} \quad (4)$$

where we exploited the orthogonality of $\mathbf{U}$. Recalling the fact that $\mathbf{\Lambda}$ is a diagonal matrix, $V[\mathbf{F}] = \mathbf{\Lambda}$ implies that the principal components are mutually uncorrelated, and have a variance equal to the associated eigenvalue. Therefore, the proportion of variance in the original n variables explained by the k -th principal component is given by:

$$\frac{\lambda_k}{\sum_{i=1}^n \lambda_i}, \quad (5)$$

and the cumulative proportion of explained variance by the first K components by:

$$\frac{\sum_{k=1}^K \lambda_k}{\sum_{i=1}^n \lambda_i}. \quad (6)$$

Mirroring the approach used by Guidolin and Pedio [2], we forecast the returns of each commodity j , $R_{t+1,j}$, with the principal component analysis to summarize a large set of macroeconomic factors:

$$R_{t+1,j} = \underbrace{\alpha_j}_{\text{intercept}} + \underbrace{\sum_{i=1}^N \beta_{ij} PC_{t,ij}}_{\text{macroeconomic}} + \underbrace{D_{t,j} \left(\sum_{k=1}^K \gamma_{kj} C_{t,kj} \right)}_{\text{commodity-specific}} + \underbrace{\varepsilon_{t,j}}_{\text{error term}} \quad (7)$$

where α_j is the model's intercept,

4.2 Extension I: Penalized Regression

In the first alternate approach, we keep the principal component analysis and still extract ten principal components from the n initial variables, but instead of running stepwise AIC on those selected principal components, we run a LASSO regression as an alternative way to further reduce the number of parameters to be estimated. With this only being a narrow replacement compared to the original framework, every difference in results should be attributable solely to the selection method.

$$R_{t+1,j} = \alpha_j + \sum_{i=1}^{10} \beta_{ij}^{\text{LASSO}} PC_{t,ij} + \sum_{k=1}^K \gamma_{kj}^{\text{LASSO}} C_{t,kj} + \varepsilon_{t,j} \quad (8)$$

A broader extension to the approach by Guidolin and Pedio [2] is foregoing the principal component analysis entirely and applying LASSO directly to the n variables.

$$R_{t+1,j} = \alpha_j + \sum_{i=1}^{128} \beta_{ij}^{\text{LASSO}} X_{t,i} + \sum_{k=1}^K \gamma_{kj}^{\text{LASSO}} C_{t,kj} + \varepsilon_{t,j} \quad (9)$$

4.3 Extension II: Individual Fundamental Predictors

4.4 OOS Evaluation Criteria

4.5 Portfolio Back-Testing Framework

5 Data

5.1 Commodity Futures Returns

5.2 Macroeconomic Factors

5.3 Aggregate Commodity-Specific Factors

5.4 Individual Fundamental Predictors

5.5 Sample Splits and Descriptive Statistics

6 Replication & Baseline Results

6.1 In-Sample Estimation (1989-2003)

6.2 OOS Predictive Accuracy (2004-2012)

6.3 Portfolio Back-Testing (2004-2012)

6.4 Replication Assessment

7 Extension I: ML Variable Selection

7.1 Model Specification and Implementation

7.2 Variable Selection Patterns

7.3 OOS Forecasting Results - Original Sample (2004-2012)

7.4 OOS Forecasting Results - Extended Sample (2013-2024)

7.5 Portfolio Back-Testing

7.6 Robustness Checks

8 Extension II: Individual Commodity Fundamentals

8.1 Predictor Construction by Commodity Group

8.2 Individual vs. Aggregate Factors: Are They Substitutes?

8.3 OOS Forecasting Results (2004-2024)

8.4 Combined Model

8.5 Portfolio Back-Testing

8.6 Regime Interaction

9 Synthesis & Discussion

9.1 Model Horse Race

9.2 Statistical vs. Economic Loss Functions Revisited

9.3 Implications for Commodity Market Segmentation

9.4 Practical Implications

10 Conclusions

10.1 Summary of Findings

10.2 Limitations

10.3 Future Research Directions

References

- [1] Marta Giampietro, Massimo Guidolin, and Manuela Pedio. “Estimating Stochastic Discount Factor Models with Hidden Regimes: Applications to Commodity Pricing”. In: *European Journal of Operational Research* 265.2 (Mar. 1, 2018), pp. 685–702. ISSN: 0377-2217. DOI: 10.1016/j.ejor.2017.07.045.
- [2] Massimo Guidolin and Manuela Pedio. “Forecasting Commodity Futures Returns with Stepwise Regressions: Do Commodity-Specific Factors Help?” In: *Annals of Operations Research* 299.1–2 (Apr. 2021), pp. 1317–1356. ISSN: 0254-5330, 1572-9338. DOI: 10.1007/s10479-020-03515-w.
- [3] Tianxiang Hao et al. “Predicting Commodity Returns Through Image-Based Price Patterns”. In: *Journal of Futures Markets* 45.12 (2025), pp. 2434–2456. ISSN: 1096-9934. DOI: 10.1002/fut.70043.
- [4] Duc Khuong Nguyen and Thomas Walther. *Modeling and Forecasting Commodity Market Volatility with Long-term Economic and Financial Variables*. Dec. 3, 2018. DOI: 10.2139/ssrn.3294967. Social Science Research Network: 3294967. URL: <https://papers.ssrn.com/abstract=3294967> (visited on 06/14/2026). Pre-published.
- [5] Sinno Jialin Pan and Qiang Yang. “A Survey on Transfer Learning”. In: *IEEE Transactions on Knowledge and Data Engineering* 22.10 (Oct. 2010), pp. 1345–1359. ISSN: 1558-2191. DOI: 10.1109/TKDE.2009.191.
- [6] Purna Chander Mashetty et al. “Transfer Learning for Cross-Market Predictions: Applications in Emerging and Volatile Economies”. In: *2025 6th International Conference on Data Intelligence and Cognitive Informatics (ICDICI)*. 2025 6th International Conference on Data Intelligence and Cognitive Informatics (ICDICI). July 2025, pp. 621–626. DOI: 10.1109/ICDICI66477.2025.11135015.
- [7] Benjamin R. Auer. “Have Trend-Following Signals in Commodity Futures Markets Become Less Reliable in Recent Years?” In: *Financial Markets and Portfolio Management* 35.4 (Dec. 1, 2021), pp. 533–553. ISSN: 2373-8529. DOI: 10.1007/s11408-021-00385-5.
- [8] Massimo Guidolin and Serena Ionta. “Predicting Commodity Returns with Climate Variables: Statistical Loss Functions vs. Economic Value”. In: *Economics Letters* 265 (June 1, 2026), p. 113028. ISSN: 0165-1765. DOI: 10.1016/j.econlet.2026.113028.
- [9] Mostafa Moghadam. *Forecasting Asset Returns with Sentiment-Enhanced FIGARCH Models*. Dec. 23, 2025. DOI: 10.2139/ssrn.5943491. Social Science Research Network: 5943491. URL: <https://papers.ssrn.com/abstract=5943491> (visited on 06/13/2026). Pre-published.

- [10] Kais Tissaoui and Nadia Azibi. “Sentiment Momentum and Regime-Dependent Predictability in Crude Oil Markets: A Markov Switching Analysis with Higher-Order VIX and MOVE Derivatives”. In: *International Review of Economics and Finance* 108 (2026). ISSN: 1059-0560. DOI: 10.1016/j.iref.2026.105197.
- [11] Hao Yang, Jie Yang, and Yun Feng. “Climate Physical Risks and the Vulnerability of Global Agricultural Commodities”. In: *Economics Letters* 258 (Jan. 1, 2026), p. 112748. ISSN: 0165-1765. DOI: 10.1016/j.econlet.2025.112748.

*not cited yet

List of Acronyms

CNN Convolutional Neural Network

MAE Mean Absolute Error

RMSE Root Mean Squared Error