

Transfer Learning for Cross-Market Commodity Return Forecasting

Leo Specks
Supervisor: Prof. Massimo Guidolin
MSc Finance, Università Bocconi

July 2027

Last Modified: Monday 7th September, 2026 (18:41)
Download latest version [here](#).

Contents

1	Introduction	4
1.1	Motivation & Background	4
1.2	The Research Gap	4
1.3	Research Questions & Hypotheses	4
1.4	Contributions to the Literature	4
1.5	Thesis Outline	4
2	Literature Review	5
2.1	Commodity Forecasting	5
2.1.1	Classical Approaches	5
2.1.2	Machine Learning & Deep Learning in Commodity Markets	6
2.2	Data Scarcity in Emerging Economies	7
2.3	Cross-Market Commodity Dynamics	7
2.4	Transfer Learning: Foundations & Applications	7
2.4.1	What is Transfer Learning?	7
2.4.2	Transfer Learning in Finance	7
2.4.3	Domain Adaptation Techniques	7
2.5	Research Gap & Positioning	7
3	Methodology	8
3.1	Framework Overview	8
3.1.1	Source and target domains	8
3.1.2	The comparison ladder	9
3.1.3	The three-stage pipeline	11
3.1.4	Information set and temporal discipline	12
3.2	Stage A: Source Model Pre-Training	13
3.2.1	Model class and parameter budget	13
3.2.2	Input representation	14
3.2.3	Normalisation protocol	14
3.2.4	Prediction target and task loss	15
3.2.5	Regularisation and validation during pre-training	16
3.2.6	What Stage A is intended to capture	16
3.3	Stage B: Domain Adaptation	17
3.3.1	Motivation and formal statement of domain shift	17
3.3.2	What is aligned, and with which data	18
3.3.3	Maximum Mean Discrepancy (MMD)	18
3.3.4	Domain-Adversarial Neural Network (DANN)	19
3.3.5	Assumptions and limitations	20
3.4	Stage C: Fine-Tuning & Target Prediction	21
3.4.1	Fine-tuning strategy	21
3.4.2	Fine-tuning objective and optimisation	22
3.4.3	Regularisation, validation and the risk of alignment drift	23
3.5	Baseline Models	23
3.5.1	Tier A: naive and simple benchmarks	24
3.5.2	Random Forest (supplementary)	24

1	3.5.3	Tier B: target only-neutral model	25
2	3.6	Evaluation Metrics	25
3	3.6.1	Point-forecast accuracy	26
4	3.6.2	Directional accuracy	26
5	3.6.3	Statistical comparison	27
6	3.6.4	Economic-significance metric	27
7	3.7	Hyperparameter Selection & Reproducibility	28
8	3.7.1	The target validation resource: a rolling-origin scheme	28
9	3.7.2	Architecture capacity	28
10	3.7.3	Tier-specific hyperparameter grids	28
11	3.7.4	Freezing depth: a single default, not a scenario-dependent grid	28
12	3.7.5	Re-estimation schedule for rolling evaluation	28
13	3.7.6	Seeding and stochasticity	28
14	3.7.7	Documentation	28
15	4	Data	29
16	4.1	Commodity Classification & Selection	29
17	4.2	Source Market: Global Benchmark Exchanges	29
18	4.3	Target Markets: Data-Sparse Economies	30
19	4.4	Feature Engineering	30
20	4.5	Descriptive Statistics & Stylised Facts	30
21	4.6	Train/Validation/Test Split	30
22	5	Empirical Results	31
23	5.1	Main Results: Prediction Performance	31
24	5.2	Ablation Study: Decomposing Transfer Learning Gains	31
25	5.3	Cross-Country Variation	31
26	5.4	Cross-Commodity Variation	31
27	5.5	Robustness Checks	31
28	5.5.1	Alternative Architectures	31
29	5.5.2	Alternative Source Markets	31
30	5.5.3	Sub-sample Stability	31
31	5.5.4	DANN vs. MMD	31
32	5.6	Economic Significance: Trading Strategy Simulation	31
33	6	Discussion	32
34	6.1	Interpretation of Core Findings	32
35	6.2	Comparison with Related Work	32
36	6.3	Economic and Policy Implications	32
37	6.4	Limitations	32
38	6.5	Future Research Directions	32
39	7	Conclusions	33
40	7.1	Summary of the Thesis	33
41	7.2	Answer to the Research Question	33
42	7.3	Broader Significance	33
43	7.4	Closing Remarks	33

1	References	34
2	List of Acronyms	48
3	Appendix	49

1 Introduction

1.1 Motivation & Background

Some of the more complex models seen in Section 2.1.1 suffer from needing of a lot of data. For some more niche markets with shorter histories and more constrained datasets this raises a problem. The most straightforward and perhaps naive solution is to simply train models on data-rich markets and implement them on data-sparse ones. However, it has been shown that traditional models fail when being transferred cross-market. When trying to transfer a convolutional neural network (CNN) model trained on US commodity futures to predict returns of the Chinese market, Hao et al. [1] find that this model fails to systematically outperform models trained on data specific to the Chinese market. These results suggest that there are little significant return-predictive patterns from US commodity futures which generalize across different markets. Hao et al. [1] give two explanations on why employing models cross-market fails in commodity markets. **insert reasons** Furthermore, Nguyen and Walther [2] have also found results to be heterogeneous across different commodity futures as no single model fit all commodities. It seems that traditional models lack cross-market transferability across both different countries and commodities.

1.2 The Research Gap

1.3 Research Questions & Hypotheses

1.4 Contributions to the Literature

1.5 Thesis Outline

2 Literature Review

This chapter will review the existing literature on the topic and work out the research gap this thesis aims to explore. Section 2.1 goes over existing approaches used in commodity forecasting and presents previous and remaining challenges. Section 2.4 explores the field of transfer learning as a potential solution for some of these remaining challenges. Section 2.5 shows where the existing literature leaves a gap that this thesis aims to fill, as the intersection of transfer learning, commodity markets and data-sparse emerging economies.

2.1 Commodity Forecasting

Before evaluating whether transfer learning improves commodity forecasting in data-sparse markets, it is necessary to establish what classical, single-market methods already achieve when data is not scarce. This benchmark matters for two reasons: classical models supply the target-only baselines against which the proposed framework is compared, see Section 3.5, and their known limitations, even under the most favourable data conditions, motivate the case for knowledge transfer in the first place.

2.1.1 Classical Approaches

A large body of empirical literature deals with regressions of commodity series on macroeconomic and financial predictors and evaluates the in-sample and out-of-sample forecasting performance. For instance, Gargano and Timmermann [3] study the predictive power over different forecast horizons, economic states, and types of commodities. Using Reuters/Jefferies-CRB commodity indexes over 1947-2010, they find that predictability from variables such as industrial production growth, commodity-currency exchange rates and global economic activity indices is present but varies highly between different commodities. While it is strongest for industrials, metals and the broad commodity index, and weaker for fats/oils, foods and livestock. Furthermore, they find forecasting performance to be more pronounced during recessions. Previously, Chen, Rogoff, and Rossi [4] showed that the exchange rates of only five commodity exporters (Australia, Canada, New Zealand, South Africa and Chile) have solid forecasting power over a global aggregate commodity price index made up by more than forty products traded on various exchanges and published by the IMF. Further studies find more limited results for the predictability of commodity spot indices through commodity exchange rates and other macroeconomic

variables [5]. Guidolin and Pedio [6] examine whether commodity-specific factors (basis, hedging pressure and momentum) improve forecasting performance. The idea is to test whether they include information not already captured by common macroeconomic variables. While including commodity-specific variables does not improve the forecasting performance, their importance seems to depend on whether statistical or economic loss functions are used. This apparent disconnect between statistical forecast accuracy of returns and the economic value created when included in a portfolio choice problem is also present in other studies [7, 8].

2.1.2 Machine Learning & Deep Learning in Commodity Markets

Over the past two decades commodity forecasting has shifted from classical statistical approaches toward machine learning and deep learning techniques. Traditional models such as autoregressive integrated moving average (ARIMA), generalized autoregressive conditional heteroskedasticity (GARCH) and exponential smoothing rely on assumptions of stationarity and linearity that frequently fail for commodity price series characterized by nonlinearity, volatility clustering and structural breaks [9, 10]. Machine learning methods like support vector regression (SVR) or extreme gradient boosting (XGBoost) provide greater flexibility in modeling multidimensional relationships without imposing strict distributional assumptions [11]. Deep learning architectures, particularly multilayer perceptron (MLP), recurrent neural network (RNN), long short-term memory (LSTM), gated recurrent unit (GRU), echo state model (ESN) and bidirectional long short-term memory (BiLSTM) have further advanced the field by capturing long-term temporal dependencies and complex nonlinear patterns [11, 10]. Next to individual approaches, hybrid models that combine decomposition techniques, ensemble methods and multiple neural architectures have emerged as a dominant trend. Fang et al. [12] use a combination of three models, support vector machine (SVM), neural network (NN) and ARIMA, to produce short-term forecasts and find that it outperforms the individual ones. Ray et al. [13] report similar results with a comparable ensemble model using a monthly forecasting period.

The literature spans multiple commodity sectors, with agricultural prices receiving a lot of attention due to their volatility and implications for food security [11, 14]. Energy commodities, especially crude oil, constitute a second cluster, driven by their geopolitical importance and market liquidity [9, 15, 16]. Studies like those from Varshini, Kayal, and Maiti [17] or Jin and Xu [18] cover forecasting of precious and industrial metals as another

category, with both general economic and financial importance.

2.2 Data Scarcity in Emerging Economies

2.3 Cross-Market Commodity Dynamics

2.4 Transfer Learning: Foundations & Applications

This shortcoming naturally invites the exploration of methodologies, which aim to build predictors, that allow for differences cross-market. The goal is to leverage data-rich markets, while still being able to include market-specific characteristics. This approach can be categorised as a transfer learning technique, which is a subset of machine learning, that allows for the assumption of different domains, tasks and distributions between the training and test set [19].

2.4.1 What is Transfer Learning?

2.4.2 Transfer Learning in Finance

2.4.3 Domain Adaptation Techniques

Mashetty et al. [20] explain approach This approach proved to outperform both a traditional machine learning model and a deep learning model, which were trained only on target data, in all three chosen metrics, namely mean absolute error (MAE), root mean squared error (RMSE) and directional accuracy.

2.5 Research Gap & Positioning

3 Methodology

This chapter will build the transfer learning framework used to test whether forecasting performance can be improved compared to known approaches. Section 3.1 presents the three-stage pipeline, which mirrors the methodology of Mashetty et al. [20], and explains how the stages connect and what information flows between them. Sections 3.2, 3.3 and 3.4 each explore one of the three stages: source pre-training, domain adaptation as well as fine-tuning and target prediction.

3.1 Framework Overview

This section introduces the research design aimed at answering the research questions introduced in Section 1.3. The goal throughout this thesis is to forecast the h -period-ahead log return of a commodity traded in a data-scarce target market,

$$r_{t+h} = \ln P_{t+h} - \ln P_t \quad (1)$$

where P_t denotes the commodity price at time t and h is the forecast horizon. Returns are used as the prediction target rather than price levels because commodity price series are typically non-stationary, whereas returns provide a more suitable basis for forecasting while retaining a clear economic interpretation. The choice of h depends on the length and characteristics of the available target-market sample and is justified in Section 4.6.

For each forecast origin t , the models described in the remainder of the chapter generate a forecast $\hat{r}_{t+h}$ using only information that would have been available at time t .

Figure 1 shows a summary of the framework, including the data sources, the three stages of the pipeline, the resulting comparison ladder and the common evaluation module. Each tier is evaluated using the same target test window, forecast horizon and predictor availability.


3.1.1 Source and target domains

The empirical framework separates the analysis into a source and target-domain. The source-domain contains historical observations from a relatively data-rich, developed commodity market and is represented as

$$D_S = \{(x_s^{(i)}, y_s^{(i)})\}_{i=1}^n, \quad (2)$$

Figure 1

Overview of the source-pre-training, domain-adaptation, fine-tuning pipeline and the resulting five-tier model comparison ladder



where $x_s^{(i)}$ denotes the vector of predictors available at the relevant forecast origin and $y_s^{(i)}$ is the corresponding realised return.

The target domain is defined in the same way but contains a substantially smaller sample from the data-scarce market of interest:

$$D_T = \left\{ \left(x_t^{(j)}, y_t^{(j)} \right) \right\}_{j=1}^m, \quad m \ll n. \quad (3)$$

Transfer learning is motivated precisely by the limited amount of target-domain data relative to the larger information set available in the source domain ($m \ll n$). The specific commodities included in the analysis, together with the choice of source and target markets are established in Chapter 4 according to data availability and economic relevance.

3.1.2 The comparison ladder

The main empirical question is whether models that draw on the source-domain dataset D_S achieve better forecasting performance than otherwise comparable models that rely only on target-domain information, when both are evaluated over the same target sample. To evaluate this question more granularly the empirical design distinguishes five model tiers according to the information available to each model and the way in which that information is used.

- **Tier A: Naive benchmark.** This tier consists of a simple forecasting rule based only on the target series, such as a zero-return, random-walk or historical-mean forecast. It contains no parameters estimated from the source-domain and provides a basic benchmark that any learned model should be expected to outperform.
- **Tier B: Target-only model.** Models on this tier are also trained only on the target-domain dataset D_T . However, its functional form should be identical, or as closely comparable as possible, to that of the transfer-learning models. Tier B is supposed to separate the effect of the model architecture from the effect of transfer learning. If a transfer model performs better than Tier A but not Tier B, the improvement cannot be attributed to the use of source-domain information.
- **Tier C: Direct transfer.** The source-trained model $f_S(x; \theta_S)$ is applied directly to target-domain observations without any target-specific parameter updating. This provides a measure of how well knowledge learned in the source market transfers to the target market in its original form. It also serves as the reference point for assessing the additional contribution of fine-tuning and domain adaptation in Tiers D and E.
- **Tier D: Fine-tuned transfer.** The source-trained parameters θ_S are updated through supervised training on D_T , but no separate domain-adaptation mechanism is introduced.
- **Tier E: Domain-adapted transfer.** The source-trained model undergoes an explicit domain-adaptation step before, or jointly with, supervised fine-tuning on D_T .

The resulting tier structure allows the contribution of each component to be examined separately. Comparisons of Tier D or E with Tier B assess whether the use of source-domain information adds value relative to training on the target-domain alone. The comparison between Tiers D and C isolates the effect of target-domain fine-tuning, while the comparison between Tier E and D captures the additional contribution of explicit domain adaptation. These comparisons correspond to the thesis hypotheses and are considered directly in Table 1. Accordingly, the performance of Tier E is not interpreted in isolation as evidence in favour of transfer learning. The lower tiers are reported such that any potential gains can be assessed against appropriately demanding benchmarks.

As described in Section 3.5, Tier B is estimated in two forms that differ only in the predictors available to the model. The restricted specification uses the set of features common to both the source and target domains, providing a like-for-like comparison with Tiers C-E. The full specification additionally includes predictors that are available only

in the target market. Comparing these two versions helps determine whether differences in forecasting performance arise from transferred information or from the availability of additional local predictors.

Table 1 shows how each of the thesis hypotheses corresponds to a specific comparison analyzed in Chapter 5.

Table 1

Mapping the comparison ladder to the research hypotheses

Hypothesis	Comparison	Evaluated in
H1: transfer improves on target-only training	Tier D and/or Tier E vs. Tier B (restricted)	Sections 5.1, 5.2
H2: transfer's advantage is larger when target data are scarcer	Gap between the best-performing transfer tier and Tier B, evaluated across shrinking target training histories	Section 5.3
H3: adaptation improves on naive transfer	Tier D and/or Tier E vs. Tier C	Section 5.2
H4: transfer performance depends on source-target similarity	Ranking of Tiers C-E under alternative source-market choices, related to ex-ante economic and statistical similarity measures	Section 5.5.2

3.1.3 The three-stage pipeline

Tiers C, D and E are constructed using a shared three-stage pipeline, drawing on the general source pre-training, domain-adaptation and fine-tuning framework by Mashetty et al. [20] for cross-market financial prediction.

- **Stage A: Source pre-training.** A model $f_S(x; \theta_S)$ is trained on D_S to minimise a supervised prediction loss, producing pre-trained parameters θ_S that are intended to capture return-generating patterns that are not specific to source market alone (e.g., momentum, mean reversion or responses to common global factors). Applying $f_s(\cdot; \theta_S)$ directly to D_T without further modification yields Tier C.
- **Stage B: Domain adaptation.** A separate alignment procedure is introduced to reduce differences between the source and target feature distributions. This addresses the possibility that a model trained solely to minimise prediction error in the source domain may learn representations that do not transfer effectively to the target domain. The precise alignment mechanism, such as a discrepancy-based penalty using maximum mean discrepancy (MMD) or an adversarial approach, is specified in Section 3.3 after the model class has been established.

- **Stage C: Fine-tuning.** The pre-trained, and where applicable adapted, model parameters are subsequently updated through supervised training on D_T . Fine-tuning the output of Stage A directly produces Tier D, whereas fine-tuning the model after Stage B produces Tier E.

The framework as it is used in this thesis differs from Mashetty et al. [20] in several ways due to its application to commodity rather than equity markets. First, clearly the source and target domains are defined using commodity return series rather than equity-index ones. This distinction is important because commodity-market linkages may operate through mechanisms such as global price discovery, storage and convenience-yield dynamics or commodity-specific supply shocks, which differ from the forces underlying equity-markets comovement. Second, the empirical design considers multiple target markets and, where data availability allows, multiple commodity classes. Third, the contribution of each stage of the pipeline is assessed separately through the explicit ablation structure represented by Tiers C, D and E, rather than evaluating only the fully combined model.

3.1.4 Information set and temporal discipline

Since the aim of the analysis is to perform a genuine out-of-sample forecast, the same chronological restrictions are applied throughout the pipeline. For any forecast at time t , models can only use predictor information available at or before time t . Target domain observations occurring after the forecast origin are excluded from all stages of model development and evaluation.

In Stage A, source-domain pre-training is conducted using only observations available before the beginning of the target evaluation period. This ensures that the pre-trained model is based solely on information that would have been available when the target forecasting started. If an expanding- or rolling-window retraining procedure is used rather than a single fixed pre-training step, the relevant procedure is described in Section 3.2.

Predictor transformations, such as normalisation, scaling, and feature construction, are estimated using the training data only. The fitted transformations are then applied unchanged to the validation and test sets. These chronological and preprocessing rules apply consistently across all stages and tiers, and are therefore stated here rather than repeated in each subsequent section unless a particular stage requires additional clarification.

3.2 Stage A: Source Model Pre-Training

This section defines Stage A of the pipeline introduced in Section 3.1.3, in which the source-domain model $f_S(x; \theta_S)$ is estimated using D_S . The resulting parameters θ_S are subsequently carried forward into Tiers C, D and E. Therefore, the modelling decisions made at this stage determine the information and structure available to all three transfer-learning tiers.

3.2.1 Model class and parameter budget

The underlying model $f(\cdot; \theta)$ is specified as a supervised sequence-learning architecture. Two candidate models are used: a LSTM network and a temporal convolutional network (TCN)¹. Both are compatible with the pipeline described in Section 3.1.3 and with the tiered comparison framework introduced in Section 3.1.2. The final choice between the two architectures or the decision to report results for both as part of the robustness analysis in Section 5.5 is made after the source- and target domain sample sizes have been established in Chapter 4.

The comparison design imposes an additional constraint on model selection that would not necessarily arise in a conventional pre-training and fine-tuning setting. Tier B, described in Section 3.3, is trained from scratch using only target-domain data and employs the same or a closely comparable model architecture. The underlying model can therefore not be selected solely on the basis of its performance on D_S . An architecture that fits the relatively large source-domain sample well but has too many parameters for the smaller target-domain dataset D_T could place Tier B at a structural disadvantage. In that case, superior performance by Tiers D or E might reflect differences in effective model capacity rather than the additional value of information from the transfer.

For this reason, the size of the architecture is constrained by the amount of data available to the restricted Tier B specification. Parameters such as hidden-state dimension, network depth and lookback length are chosen such that specification can be estimated without pronounced overfitting. This condition is assessed empirically through learning curves that relate out-of-sample forecasting error to the size of the target-domain training sample for a given level of model capacity, as part of the hyperparameter-selection procedure described in Section 3.7. Where this requirement substantially restricts the capacity that could otherwise be achieved in Stage A, the resulting trade-off is reported explicitly.

¹Finalize which models will be used.

3.2.2 Input representation

The input to f_S is based exclusively on a shared-core feature set which are consistent with those found in the literature². This includes own-series technical indicators, such as lagged log returns, rolling realised volatility, moving averages, the relative strength index and Bollinger-band measures, together with global and macro-financial variables that are plausibly relevant across commodity markets. Examples include a global commodity index or benchmark return, the US dollar index, a global risk-sentiment measure and global short-term interest rates. No target-specific local predictors are included at this stage. As discussed in Section 3.2.6 and formalised in Section 3.5, such variables are not available in the source domain and therefore cannot be incorporated into a source-trained model that is intended to be applied directly to D_T .

Let $x_\tau \in \mathbb{R}^d$ denote the shared-core feature vector observed at time τ , where the feature dimension d is the same in the source and target domains. Each source-domain training observation is constructed using a lookback window of length L .

$$X_s^{(i)} = (x_{i-L+1}, x_{i-L+2}, \dots, x_i), \quad y_s^{(i)} = r_{i+h}, \quad (4)$$

The source model therefore maps an $L \times d$ sequence of predictors to a scalar return forecast. The lookback length L is treated as a hyperparameter and is selected using source-domain validation data only, as described in Section 3.2.5. Its final value is reported in Section 3.7 alongside other architectural choices.

Restricting x_τ to the shared-core feature set ensures that the same input construction can be used for the target-domain observations in Tiers C, D and E. In this way, the predictor representation remains identical across domains, so differences in forecasting performance can be attributed to the modelling stage being evaluated and not to the information supplied to each model.

3.2.3 Normalisation protocol

Feature scaling is estimated once using the training partition of D_S and is then applied unchanged wherever the shared-core feature representation is used. This includes Stage A training, the zero-shot evaluation in Tier C and the domain-adaptation and fine-tuning procedures in Tiers D and E. By contrast, both the restricted and full versions of Tier B are normalised separately using statistics computed from their own target-domain training

²Insert specific citations.

partitions, since these models are estimated independently of the source-trained parameters θ_S .

The purpose of retraining the source-fitted transformation across the transfer-learning pipeline is to preserve economically meaningful differences between the source and target domains. Re-estimating the scaler on target data before the transfer model is applied would remove some of the differences in scale and volatility that form part of the domain shift being studied.

The precise scaling method is determined after examining the data in Chapter 4. The candidates are standard scaling based on the mean and variance and robust scaling based on the median and interquartile range. Given the tendency of commodity and emerging-market return series to exhibit heavy tails, the robust specification is treated as the default unless the descriptive evidence reported in Section 4.5 provides a reason to prefer an alternative.

Applying source-based scaling to target observations may nevertheless produce transformed values outside the range encountered during Stage A, particularly when volatility is substantially greater in the target market. This possibility is relevant to the negative-transfer mechanisms discussed in Chapter 6. To improve numerical stability during fine-tuning, **gradient clipping** is therefore employed, while extreme transformed inputs may be winsorised using thresholds determined again exclusively from the source training distribution. These measures are intended as stability safeguards and should not alter the underlying normalisation strategy.

3.2.4 Prediction target and task loss

Stage A is formulated as a regression task using the h -period-ahead log return defined in Section 3.1, $y = r_{r+h}$. Return direction is **not** modelled separately. Instead, directional accuracy, as defined in Section 3.6, is calculated from the sign of the continuous forecasts. This avoids introducing an additional classification component into the pipeline where it would not provide a clear identification advantage.

The source-domain is trained by minimising mean squared error:

$$\mathcal{L}_S(\theta_S) = \frac{1}{n} \sum_{i=1}^n (f_S(X_s^{(i)}; \theta_S) - y_s^{(i)})^2, \quad (5)$$

Optimisation is performed using mini-batch gradient descent with the Adam optimiser. The learning rate, batch size and number of training epochs are treated as hyperparameters.

and selected using source-domain validation data, as described in Section 3.2.5. The fine-tuning stage uses a reduced learning rate, with the adjustment procedure specified in Section 3.4.

3.2.5 Regularisation and validation during pre-training

Stage A incorporates dropout regularisation and early stopping to reduce the risk of overfitting to D_S and to encourage the model to learn representations that are less dependent on source-market-specific features. The early-stopping rule and all architectural hyperparameters are selected using time-series cross-validation conducted exclusively within the source domain. Validation folds are constructed chronologically using expanding or rolling training windows, with no reandom partitioning of observations.

In line with the chronological restrictions set out in Section 3.1.4, all source-domain training and validation observations used for these decisions precede the beginning of the target evaluation period. No target-domain data, whether from training, validation or testing, are used to inform model selection or hyperparameter tuning at this stage. The complete search procedure and the final hyperparameter values are reported in Section 3.7.

3.2.6 What Stage A is intended to capture

The extent to which the pre-trained parameters θ_S are valuable in the target domain depends on whether they capture relationships that generalise beyond the source market. In this setting, two broad groups of shared inputs are especially relevant, although their potential for transfer is likely to differ.

The first group comprises patterns associated with the underlying return-generating process. These include short-term momentum or mean reversion, volatility clustering and the sensitivity of commodity returns to global risk and financing conditions, such as movements in the US dollar index, global short-term interest rates and measures of global risk sentiment. Such relationships are plausible candidates for cross-market transfer because the mechanisms underlying them, including global commodity price discovery, dollar-denominated pricing, and common exposure to global business-cycle conditions, extend beyond any individual source market.

The second group consists of relationships that may arise from features specific to the source market itself. Examples include trading-session effects, liquidity conditions on the source exchange and market responses to source-country economic announcements. These patterns are less likely to generalise to the target domain. Their influence on θ_S is

supposed to be reduced by the domain-adaptation procedure in Stage B as described in Section 3.3. However, this distinction should be interpreted as a working hypotheses about the types of information that source pre-training may capture and not as an empirical finding established in advance. The issue is examined again in Section 5.4 through the cross-commodity comparison and in Chapter 6 in the discussion of negative transfer.

3.3 Stage B: Domain Adaptation

This section defines Stage B of the pipeline, which introduces the explicit domain-alignment procedure used in Tier E. Stage B begins with the pre-trained parameters θ_S and the source-fitted feature representation established in Section 3.2. Its purpose is to address the possibility that, even though the source and target domains use the same shared-core feature construction, the representation learned during source-domain pre-training may not transfer effectively to the target domain.

3.3.1 Motivation and formal statement of domain shift

Tier C evaluates the source-trained model $f_S(\cdot; \theta_S)$ on target-domain inputs without any parameter adjustment. Its performance relative to Tiers A and B, reported in Section 5.2, therefore provides a direct measure of how much predictive information transfers from the source market in absence of adaptation.

Zero-shot transfer may perform poorly even when the source and target domains are constructed from an identical shared-core feature set. The reason is that identical feature definitions do not imply identical feature distributions or internal model representations. Let $\phi(x; \theta_\phi)$ denote the output of the feature-extraction component of the network, consisting of all layers of f_S up to, but excluding, the final regression head, where $\theta_\phi \subset \theta_S$. The relevant form of domain shift arises when the distributions of these learned representations differ between the source and target domains:

$$P_S(\phi(X_s)) \neq P_T(\phi(X_t)), \quad (6)$$

Such differences can arise for several plausible reasons. For example, the target market may exhibit systematically higher volatility, global variables in the shared-core feature set may affect local returns differently across markets or the target economy may experience structural breaks that are absent from the source sample. A regression head calibrated using source-domain representations may therefore encounter target-domain representations that

differ substantially from those observed during pre-training. In such circumstances, reliable target predictions are not guaranteed, even if the underlying predictive relationship has not itself changed. Section 3.3.5 returns to the distinction between shifts in the distribution of model inputs and changes in the underlying predictive relationship.

3.3.2 What is aligned, and with which data

Stage B seeks to align the source and target distributions in the learned representation space $\phi(x)$ and not the level of the raw inputs, because it is this representation that is passed to the final regression head. The alignment procedure draws on two data sources: the labelled source-domain training set D_S^{train} , which allows the predictive objective to remain active while the representation is adjusted, and unlabelled target target-domain inputs $\{x_t^j\}$ taken exclusively from the target training partition D_T^{train} .

Restricting the target inputs to D_T^{train} is an important feature of the empirical design. Methods such as MMD and domain-adversarial neural network (DANN) do not require target labels and many applications of domain adaptation therefore allow access to unlabelled observations from the evaluation sample. That convention is not needed here. In a forecasting framework governed by a fixed evaluation origin, using predictor information from the target validation or test period to shape the feature representation would allow future distributional information to influence the model before evaluation, even if no target outcomes were observed.

The chronological restrictions introduced in Section 3.1.4 are therefore extended explicitly to unlabelled target data. Any target observations entering the domain adaptation loss $\mathcal{L}_{DA}$ are drawn only from D_T^{train} . The target validation and test partitions are excluded entirely from Stage B.

3.3.3 Maximum Mean Discrepancy (MMD)

The principal domain-adaptation method used in Stage B is a MMD penalty defined in a reproducing kernel Hilbert space (RKHS) $\mathcal{H}$ with associated kernel $k(\cdot, \cdot)$. Let $\phi(x_s^{(i)}), i = 1, \dots, n'$ and $\phi(x_t^{(j)}), j = 1, \dots, m'$ denote source- and target-domain representations drawn from their respective training partitions. The empirical squared MMD is given by:

$$\overline{\text{MMD}}^2 = \left\| \frac{1}{n'} \sum_{i=1}^{n'} \varphi(\phi(x_s^{(i)})) - \frac{1}{m'} \sum_{j=1}^{m'} \varphi(\phi(x_t^{(j)})) \right\|_{\mathcal{H}}^2, \quad (7)$$

where $\varphi(\cdot)$ denotes the implicit feature map associated with the kernel k . In practice,

the RKHS representation does not need to be computed explicitly. The square distance can instead be expressed entirely through kernel evaluations between pairs of source representations, pairs of target representations and source-target representation pairs. This allows the discrepancy between the two domains to be estimated directly from the learned representations.

A Gaussian radial basis function (RBF) kernel is used as the default specification. A multi-kernel formulation, following the broader approach to deep domain adaptation proposed by Long et al. [21], is retained as a possible robustness extension.

Defining the domain-adaptation loss $\mathcal{L}_{DA}$ as $\widehat{\text{MMD}}^2$, the Stage B objective becomes

$$\mathcal{L}_{\text{Total}}^{(B)}(\theta) = \mathcal{L}_S(\theta) + \lambda \mathcal{L}_{DA}(\theta_\phi), \quad (8)$$

where $\mathcal{L}_S$ is the source-domain prediction loss introduced in Section 3.2.4 and $\lambda \geq 0$ controls the balance between preserving source predictive performance and reducing the discrepancy between source and target representations. Gradients from both terms update the feature-extractor parameters θ_ϕ , while the regression head continues to receive gradient information only through $\mathcal{L}_S$, since target labels are not used at this stage.

This specification intentionally excludes any supervised target-domain loss from Stage B. It therefore differs from Mashetty et al. [20], whose treatment of domain adaptation incorporates a target prediction loss in a way that is not clearly separated from the paper’s subsequent fine-tuning stage. Using the objective $\mathcal{L}_S + \lambda \mathcal{L}_{DA}$ preserves a clear distinction between the adaptation and fine-tuning in the empirical design. This separation is important for the ablation analysis in Section 5.2: Stage B captures the contribution of representational alignment, whereas Stage C, described in Section 3.4, is the first stage at which target-domain labels are introduced.

The weighting parameter λ is selected using the labelled target validation partition, which is kept separate from both D_T^{train} and the target test set. The full procedure is described in Section 3.7 and follows the three-way target-domain split established in Section 4.6.

3.3.4 Domain-Adversarial Neural Network (DANN)

As an alternative to MMD, the thesis uses a DANN formulation based on Ganin et al. [22]. Under this approach, a domain discriminator $D(\phi(x))$ is trained to distinguish whether a learned representation originates from the source or target domain. A gradient-reversal

layer simultaneously updates the feature extractor so that the resulting representations become increasingly difficult for the discriminator to classify by domain. The corresponding adversarial objective is

$$\min_{\theta_\phi} \max_{\theta_D} \mathbb{E}_{x_s \sim D_S^{\text{train}}} [\log D(\phi(x_s))] + \mathbb{E}_{x_t \sim D_T^{\text{train}}} [\log (1 - D(\phi(x_t)))], \quad (9)$$

As with mit MMD specification, only observations from the source and target training partitions enter this objective, in accordance with Section 3.1.4.

DANN is treated as a secondary adaptation method. Adversarial optimisation can be more difficult to stabilise and requires sufficient data for the domain discriminator to learn a reliable distinction between source and target representations. This is a potentially important limitation in the present setting, where target-domain data scarcity is precisely the point of the research problem. For this reason, the principal Tier E results reported in Chapter 5 use MMD as the domain-adaptation mechanism. DANN is examined separately as part of the robustness analysis in Section 5.5.4, where its performance can be compared directly with the MMD specification without assuming that the two alignment approaches are empirically interchangeable.

3.3.5 Assumptions and limitations

Both MMD and DANN, as formulated above, seek to reduce differences in the marginal distribution of the learned $\phi(x)$ across the source and target domains. Because neither method uses target-domain labels during Stage B, neither can establish whether the conditional relationship between the representations and the forecasting target $P(y | \phi(x))$ is stable across domains. The effectiveness of representation alignment therefore depends on an assumption that the predictive relationship is sufficiently similar after accounting for differences in feature distribution.

There are economically plausible reasons why this assumption may fail. For example, a common global risk indicator may have a different relationship with returns in the target market because of differences in exchange-rate regimes, commodity-pricing structures or local market characteristics³. In such a case bringing the distributions of $\phi(x_s)$ and $\phi(x_t)$ closer together would not eliminate the underlying difference in the relationship between predictors and returns. Perhaps, greater distributional alignment may reduce the usefulness of the representation for target forecasting even while making the source and

³Potentially refer back to literature section on cross-market dynamics.

target representations more difficult to distinguish. This provides one possible mechanism through which domain adaptation could generate, instead of mitigate, negative transfer. The issue is revisited in Chapter 6 after the performance of Tier E relative to Tiers B-D has been established empirically.

The absence of target labels in Stage B also means that the adaptation procedure cannot determine independently whether improved alignment translates into better target-domain forecasts. Evidence on the value of adaptation becomes available only after target-domain fine-tuning in Stage C, described in Section 3.4, and subsequent out-of-sample evaluation in Chapter 5.

3.4 Stage C: Fine-Tuning & Target Prediction

This section defines Stage C, where labelled target-domain observations are introduced into the transfer-learning pipeline for the first time. Tier D enters this stage directly from the source-pretrained model produced in Stage A, whereas Tier E begins from the domain-adapted representation obtained in Stage B. In neither case has the regression head been trained on target-domain labels before this stage. Stage C is therefore the first stage at which labelled observations from D_T are used to update model parameters. Both tiers are fine-tuned using the same procedure described below. Holding the fine-tuning protocol constant ensures that any subsequent difference in forecasting performance between Tiers D and E can be attributed to the inclusion of the domain-adaptation stage, rather than to differences in how the models are trained on target-domain data.

3.4.1 Fine-tuning strategy

Instead of treating full fine-tuning, partial fine-tuning and layer freezing as separate procedures, this thesis represents them through a single design choice: the subset of parameters updated during Stage C, denoted by $\theta_{\text{trainable}} \subseteq \theta$. Under full fine-tuning, $\theta_{\text{trainable}} = \theta$, so all network parameters are updated.

This choice aims at balancing target-domain flexibility and the preservation of the representation learned earlier in the pipeline. Allowing a larger share of the network to update gives Stage C greater capacity to adapt to the target forecasting task, but it also increases the risk that information acquired during source pre-training or domain adaptation is overwritten by a comparatively small target training sample.⁴

⁴Remember to insert the paper where they did this.

Because the appropriate degree of freezing is likely to depend on the amount of target-domain training data available, freezing depth is selected from a small, pre-specified set of configurations and may vary across the data-scarcity scenarios examined in Section 5.3. Limiting the search to a smaller number of economically and statistically motivated configurations is consistent with the overall objective of avoiding a too open-ended model-selection exercise.

3.4.2 Fine-tuning objective and optimisation

Stage C is trained using a supervised objective defined exclusively on the target-domain training partition:

$$\mathcal{L}_T(\theta) = \frac{1}{m'} \sum_{j \in D_T^{\text{train}}} \left(f(X_t^{(j)}; \theta) - y_t^{(j)} \right)^2. \quad (10)$$

Parameter updates are restricted to the trainable subset $\theta_{\text{trainable}}$:

$$\theta_{\text{trainable}} \leftarrow \theta_{\text{trainable}} - \eta_{FT} \nabla_{\theta_{\text{trainable}}} \mathcal{L}_T(\theta). \quad (11)$$

Optimisation is performed using the adaptive moment estimation (Adam) optimiser, with a fine-tuning learning rate η_{FT} chosen to be lower than the rate used during Stages A and B. This reflects the purpose of Stage C as a controlled adjustment of an already informative parameter configuration.

No domain-adaptation penalty is retained during Stage C for either Tier D or Tier E. An alternative specification would continue to optimise a combined objective such as $\mathcal{L}_T + \lambda' \mathcal{L}_{DA}$ during Tier E's fine tuning. Such an approach is closer to the formulation used by Mashetty et al. [20] and to some discrepancy-based domain-adaptation frameworks, including the deep adaptation network approach of Long et al. [21], in which an alignment penalty remains active during supervised training. However, for the main empirical design the alignment term is restricted to Stage B. Stage C therefore applies exactly the same supervised fine-tuning procedure to Tiers D and E, with the two tiers only being different in their initial parameter values. This aims at preserving the interpretation of Tier D-Tier E comparison in Section 5.2 as a test of the incremental contribution of Stage B. If the domain-adaptation penalty remained active during Tier E's fine-tuning, any difference between the two tiers could reflect both their different initialisations and their different training objectives.

3.4.3 Regularisation, validation and the risk of alignment drift

Three design choices determine the extend to which Stage C can depart from its initial parameter configuration: the stopping point, the freezing depth and the fine-tuning learning rate η_{FT} . All three are selected using the same labelled target-domain validation partition employed for choosing λ in Section 3.3.3. Stage B and C therefore share a common validation resource.⁵ The detailed selection procedure, including candidate grid, stopping rule and tuning budget, is set out in Section 3.7.

For Tier E, supervised fine-tuning introduces an additional consideration. Because the Stage C objective contains no alignment term, there is nothing explicitly preventing the representation obtained in Stage B from changing during subsequent target-domain training. Fine-tuning on a small target sample could therefore weaken or partially reverse the alignment achieved during Stage B, reducing any advantage associated with the adapted initialisation. Whether such representational drift occurs is treated as an empirical question. After Tier E completes Stage C, the MMD between its fine-tuned target representation and the corresponding source representation is calculated using the estimator defined in Section 3.3.3. This value is compared with the MMD observed immediately after Stage B and with the corresponding measure for Tier D after fine-tuning. The outcome is reported alongside the principal Chapter 5 results and revisited in Chapter 6 when assessing possible sources of negative transfer. If the outcome indicates that Tier E's representation moves back toward the distributional discrepancy observed in Tier D, while Tier E exhibits no corresponding improvement in forecasting accuracy, this would provide a reason to examine the joint-adaptation specification discussed in Section 3.4.2. That alternative would then be introduced as a targeted robustness check in Section 5.5.

3.5 Baseline Models

This section defines the two comparison tiers that rely exclusively on target-domain information and make no use of the source domain. Tier A assesses whether the use of a neural sequence model is justified relative to a simple forecasting benchmark. Tier B evaluates, conditional on that modelling architecture, whether source pre-training and subsequent adaptation provide additional predictive value. Under a more restricted configuration, only the regression head and the uppermost feature-extraction block are trainable, while lower layers remain fixed at the values obtained from Stage A for Tier D

⁵Keep in mind and ask if it raises issues.

or Stage B for Tier E.

3.5.1 Tier A: naive and simple benchmarks

Tier A consists of three established non-neural forecasting methods. The first assumes that the underlying commodity price series is a driftless random-walk. The h -step ahead return forecast at time t is therefore given by:

$$\hat{r}_{t+h}^{RW} = 0. \quad (12)$$

The second is an expanding historical-mean forecast,

$$\hat{r}_{t+h}^{HM} = \frac{1}{t} \sum_{r \leq t} r_r. \quad (13)$$

The third is an ARIMA(p,d,q)-GARCH(1,1) model fitted to the target return series. The ARIMA component models the conditional mean, while the GARCH variance equation,

$$\sigma_t^2 = \omega + \alpha \varepsilon_{t-1}^2 + \beta \omega_{t-1}^2 \quad (14)$$

captures time variation in conditional variance. The ARIMA order (p,d,q) is selected **information criteria** based exclusively on target-domain training data. As the forecast origin advances, the model is re-estimated using an expanding window. The GARCH conditional-variance forecast is retained for the trading-strategy analysis in Section 5.6 but is not used in the forecast-accuracy comparisons considered here.

All three Tier A models rely only on the target return series and do not use the shared-core predictors introduced in Section 3.2.2. Tier A represents forecasting performance attainable under minimal information requirements and without engineered predictor variables.

3.5.2 Random Forest (supplementary)

A Random Forest model trained exclusively on target-domain data is included as a supplementary benchmark for assessing whether the additional complexity of the neural sequence architecture is justified. The Random Forest uses the shared-core feature vector x_t observed at a single time point instead of the full L -period input sequence $X_T^{(j)}$ used by the neural models. This reflects the fact that a conventional tree-based model does not directly process sequential inputs. However, several variables in the shared-core feature set,

including lagged returns and rolling indicators, retain information about recent temporal dynamics. The model is estimated only under the restricted predictor specification, making it directly comparable to Tier B-restricted rather than to Tier B-full. The Random Forest hyperparameters, including number of trees, maximum tree depth and minimum leaf size, are selected using the same target-domain validation partition used elsewhere in the chapter.

The model is not incorporated into the core comparison ladder and hypothesis mapping of Section 3.1.2. Its purpose is instead to address a separate modelling question: whether the use of neural sequence architecture is justified before considering the main question of whether transfer learning improves performance conditional of that architecture.

3.5.3 Tier B: target only-neutral model

Tier B employs the same architecture family and hidden-layer configuration used in Stage A, as described in Section 3.2.1, but is initialised randomly rather than from the source-pretrained parameters θ_S . The model is trained directly on the target-domain training partition by minimising the target loss $\mathcal{L}_T$.

The restricted and full versions of Tier B differ only in the width of the input layer required to accommodate their respective predictor sets. Tier B-restricted uses the shared-core features, while Tier B-full additionally incorporates the target-specific local predictors. The hidden-layer dimensions and network depth are otherwise held consistent across the two specifications. The freezing-depth procedure described in Section 3.4 is not relevant here because there is no pretrained representation to preserve. Both variants are normalised using scaling parameters estimated from their own target-domain training data, following Section 3.2.3 and are tuned using the same target validation partition used for selecting λ in Section 3.3.3 and the Stage C stopping criteria in Section 3.4.3.

3.6 Evaluation Metrics

This section defines the procedure used to evaluate forecasting performance after all training, domain adaptation and fine-tuning stages have been completed. It is also the first stage of the empirical framework in which the target test partition D_T^{test} is accessed. The test sample is reserved solely for final out-of-sample evaluation. Accordingly, no performance metric, decision threshold or evaluation procedure introduced in this section may influence model specification, hyperparameter tuning, early stopping or the selection

of freezing depth described earlier in the chapter.

3.6.1 Point-forecast accuracy

For each model, forecast errors over the target test window are defined as $e_t = r_{t+h} - \hat{r}_{t+h}$ over the test window, $t \in T_{\text{test}}$, where $N = |T_{\text{test}}|$ denotes the number of test observations. Point-forecast accuracy is evaluated using MAE and RMSE:

$$\text{MAE} = \frac{1}{N} \sum_{t \in T_{\text{test}}} |e_t|, \quad \text{RMSE} = \sqrt{\frac{1}{N} \sum_{t \in T_{\text{test}}} e_t^2}. \quad (15)$$

Out-of-sample R^2 is defined relative to the historical-mean benchmark introduced in Tier A instead of as a stand-alone goodness-of-fit measure⁶:

$$R_{\text{OOS}}^2 = 1 - \frac{\sum_{t \in T_{\text{test}}} e_t^2}{\sum_{t \in T_{\text{test}}} (r_{t+h} - \hat{r}_{t+h}^{HM})^2}. \quad (16)$$

Here $\hat{r}_{t+h}^{HM}$ denotes the historical-mean forecast defined in Section 3.5.1. This benchmark-relative formulation follows the standard logic of out-of-sample return-predictability evaluation. It also gives R_{OOS}^2 a direct interpretation: a positive value indicates that the model produces a lower squared forecast error than the historical-mean benchmark over the same test period.

3.6.2 Directional accuracy

Directional accuracy is defined as

$$\text{DA} = \frac{1}{N} \sum_{t \in T_{\text{test}}} 1 [\text{sign}(\hat{r}_{t+h}) = \text{sign}(r_{t+h})]. \quad (17)$$

although directional accuracy is often compared with a 50% benchmark, that threshold only makes sense when positive and negative returns occur with equal frequency. The analysis instead uses the directional performance of Tier A's historical-mean forecast as the relevant benchmark. Therefore, DA^{HM} is calculated using (17) and each model's directional accuracy is assessed relative to it.

The driftless random-walk benchmark produces a forecast of exactly zero and is therefore not included in the directional accuracy comparison.

⁶Check again if makes sense and find the paper again + cite.

3.6.3 Statistical comparison

For two competing models, the loss differential at time t is $d_t = g(e_{1,t}) - g(e_{2,t})$, where $g(\cdot)$ is the squared- or absolute error loss, the Diebold-Mariano statistic is then defined as

$$DM = \bar{d} / \sqrt{\text{Var}(\bar{d})/N}, \quad (18)$$

with the variance term estimated using a Newey-West long-run variance estimator rather than the ordinary sample variance [23]. This adjustment is required because multi-step-ahead forecast errors are generally serially correlated when $h > 1$. The truncation lag is chosen with reference to the forecast horizon h ⁷.

Given the number of models included in the empirical design, testing each one pairwise would create a substantial multiple-comparison problem. The main Diebold-Mariano tests are therefore restricted to the comparisons that correspond directly to the hypotheses mapped in Section 3.1.2. Broader comparisons across the full set of model specifications are handled using a Model Confidence Set procedure and reported in Appendix 7.4.

The interpretation of these tests also distinguishes between statistical and substantive importance. A lower RMSE or MAE does not by itself establish that two models differ significantly, while statistical significance alone does not establish that the improvement is economically meaningful. Chapter 5 therefore evaluates comparisons along three dimensions: statistical significance, based on the Diebold Mariano or Model Confidence Set results; economic magnitude, reflected in the size of the differences in MAE, RMSE or directional accuracy; and practical forecasting relevance, considered in Section 3.6.4 through the implications for trading or hedging decisions.

3.6.4 Economic-significance metric

A simple trading strategy is constructed from the sign of each model's return forecast. The resulting strategy return is defined as

$$s_t = \text{sign}(\hat{r}_{t+h}) r_{t+h} - c \cdot 1[\text{position changed at } t], \quad (19)$$

where c denotes the assumed transaction cost incurred whenever the trading position changes. The value of c is stated explicitly whenever the corresponding performance measure is reported. The economic performance is summarised using the annualised

⁷Finalise with Diebold and Mariano [23].

Sharpe ratio $\bar{s}/\hat{\sigma}(s)$, with the appropriate annualisation factor applied according to the sampling frequency.

3.7 Hyperparameter Selection & Reproducibility

This section consolidates several methodological choices identified earlier in the chapter. These include the use of the common target-domain validation set introduced in Section 3.3, the selection of the appropriate freezing depth discussed in Section 3.4 and the rolling re-estimation procedure deferred at the end of Section 3.6. It also sets out the conventions adopted for random seeding and for documenting the final model-selection and estimation procedures.

3.7.1 The target validation resource: a rolling-origin scheme

Section 4.6 partitions D_T chronologically into a development region, used for model training and validation, and a held-out test set D_T^{test} , which is not accessed until the final evaluation stage in Chapter 5. Within the development region, hyperparameters are selected using a rolling-origin validation procedure. A sequence of K chronologically ordered, non-overlapping validation windows is created by progressively expanding or rolling the training sample forward through time. Each candidate hyperparameter configuration is evaluated across all K windows and the specification with the lowest average validation loss⁸ is selected. The chosen model is then re-estimated using the full development region before forecasts are generated for D_T^{test} . This approach is used due to its advantage of not permanently allocating a portion of an already short series to validation. The validation framework is used for all target-domain hyperparameter decisions made in the chapter.

3.7.2 Architecture capacity

3.7.3 Tier-specific hyperparameter grids

3.7.4 Freezing depth: a single default, not a scenario-dependent grid

3.7.5 Re-estimation schedule for rolling evaluation

3.7.6 Seeding and stochasticity

3.7.7 Documentation

⁸Which loss function?

4 Data

This chapter will describe the empirical landscape of commodity markets and available data. Section 4.1 addresses the reasons for the specific choice of commodity classes examined. Sections 4.2 and 4.3 describe the data used for the source and target market respectively.

4.1 Commodity Classification & Selection

4.2 Source Market: Global Benchmark Exchanges


Table 2

Possible Source Data Sources/Series

Source	Freq.	Commodities	Dates
Yahoo Finance (quoted from COMEX, NY Mercantile, CBOT, CME, ICE Futures)	Daily	Gold, Silver, Copper, Crude Oil WTI, Crude Oil Brent, Natural Gas, Corn, Oats, Wheat, Rough Rice, Soybeans, Feeder Cattle, Lean Hogs, Live Cattle, Cocoa, Coffee, Cotton, Lumber, Orange Juice, Sugar	1999-2026

Figure 2

Example figure with two subfigures.



4.3 Target Markets: Data-Sparse Economies

Table 3

Possible Target Data Sources/Series

Source	Freq.	Commodities	Dates
DCE (Dalian Commodity Exchange)	Daily	No. 1 Soybean, No. 2 Soybean, Soybean Meal, Soybean Oil, RBD Palm Olein, Corn, Corn Starch, Pol RG Rice, Egg, Live Hogs, Fiberboard, Blockboard, Log, Coking Coal, Coke, Iron Ore, LLDPE, LLDPE MAF, PVC, PVC MAF, Polypropylene, Polypropylene MAF, Ethylene Glycol, Ethenylbenzene, LPG, Benzene	2006-2026
Zhengdong Commodity Exchange	Daily	Apple, Cotton, Chinese Jujube, Cotton Yarn, Flat Glass, Japonica Rice, Late Rice, Methanol, Rapeseed Oil, Polyester Staple Fiber, Wheat PM, Early Rice, Rapeseed Meal, Rapeseed, Soda Ash, Ferrosilicon, Manganese Silicon, White Sugar, PTA, Urea, Wheat WH, Thermal Coal, Peanut Kernel, Paraxylene, Sodium Hydroxide, Polyethylene Terephthalate Resin for Bottles, Propylene	2015-2026
Multi Commodity Exchange (MCX) India	Daily	Aluminium, Cardamom, Copper, Cotton, Cotton Oil, Crude Oil, ELECDMBL, Gold, Goldguinea, Goldm, Goldpetal, Goldten, Kapas, Lead, Leadmini, MCXBULLDEX, MCXMETLDEX, Metha Oil, Natural Gas Mini, Natural Gas, Nickel, Silver, Silver 100, Silver M, Silver MIC, Steel Rebar, Zinc, Zinc Mini	2003-2026

4.4 Feature Engineering

4.5 Descriptive Statistics & Stylised Facts

4.6 Train/Validation/Test Split

5 Empirical Results

5.1 Main Results: Prediction Performance

5.2 Ablation Study: Decomposing Transfer Learning Gains

5.3 Cross-Country Variation

5.4 Cross-Commodity Variation

5.5 Robustness Checks

5.5.1 Alternative Architectures

5.5.2 Alternative Source Markets

5.5.3 Sub-sample Stability

5.5.4 DANN vs. MMD

5.6 Economic Significance: Trading Strategy Simulation

6 Discussion

6.1 Interpretation of Core Findings

6.2 Comparison with Related Work

6.3 Economic and Policy Implications

6.4 Limitations

6.5 Future Research Directions

7 Conclusions

7.1 Summary of the Thesis

7.2 Answer to the Research Question

7.3 Broader Significance

7.4 Closing Remarks

References

- [1] Tianxiang Hao et al. “Predicting Commodity Returns Through Image-Based Price Patterns”. In: *Journal of Futures Markets* 45.12 (2025), pp. 2434–2456. ISSN: 1096-9934. DOI: 10.1002/fut.70043.
- [2] Duc Khuong Nguyen and Thomas Walther. *Modeling and Forecasting Commodity Market Volatility with Long-term Economic and Financial Variables*. Dec. 3, 2018. DOI: 10.2139/ssrn.3294967. Social Science Research Network: 3294967. URL: <https://papers.ssrn.com/abstract=3294967> (visited on 06/14/2026). Pre-published.
- [3] Antonio Gargano and Allan Timmermann. “Forecasting Commodity Price Indexes Using Macroeconomic and Financial Predictors”. In: *International Journal of Forecasting* 30.3 (July 1, 2014), pp. 825–843. ISSN: 0169-2070. DOI: 10.1016/j.ijforecast.2013.09.003.
- [4] Yu-Chin Chen, Kenneth S. Rogoff, and Barbara Rossi. “Can Exchange Rates Forecast Commodity Prices?” In: *The Quarterly Journal of Economics* 125.3 (Aug. 1, 2010), pp. 1145–1194. ISSN: 0033-5533. DOI: 10.1162/qjec.2010.125.3.1145.
- [5] Jan J.J. Groen and Paolo A. Pesenti. “Commodity Prices, Commodity Currencies, and Global Economic Developments”. In: *Commodity Prices and Markets*. Ed. by Takatoshi Ito and Andrew K. Rose. University of Chicago Press, Mar. 1, 2011, p. 0. ISBN: 978-0-226-38689-8. DOI: 10.7208/chicago/9780226386904.003.0002.
- [6] Massimo Guidolin and Manuela Pedio. “Forecasting Commodity Futures Returns with Stepwise Regressions: Do Commodity-Specific Factors Help?” In: *Annals of Operations Research* 299.1–2 (Apr. 2021), pp. 1317–1356. ISSN: 0254-5330, 1572-9338. DOI: 10.1007/s10479-020-03515-w.
- [7] Tolga Cenesizoglu and Allan Timmermann. “Do Return Prediction Models Add Economic Value?” In: *Journal of Banking & Finance*. International Corporate Finance Governance Conference 36.11 (Nov. 1, 2012), pp. 2974–2987. ISSN: 0378-4266. DOI: 10.1016/j.jbankfin.2012.06.008.
- [8] Massimo Guidolin and Serena Ionta. “Predicting Commodity Returns with Climate Variables: Statistical Loss Functions vs. Economic Value”. In: *Economics Letters* 265 (June 1, 2026), p. 113028. ISSN: 0165-1765. DOI: 10.1016/j.econlet.2026.113028.

- [9] P. Foroutan and Salim Lahmiri. “Deep Learning Systems for Forecasting the Prices of Crude Oil and Precious Metals”. In: *Financial Innovation* 10 (July 16, 2024). DOI: 10.1186/s40854-024-00637-z.
- [10] Kaixuan Sun, Qi Yao, and Yanhui Li. “A Novel Agricultural Commodity Price Prediction Model Integrating Deep Learning and Enhanced Swarm Intelligence Algorithm”. In: *PLOS One* 20 (Dec. 2, 2025). DOI: 10.1371/journal.pone.0337103.
- [11] R. Manogna, Vijay Dharmaji, and S. Sarang. “A Novel Hybrid Neural Network-Based Volatility Forecasting of Agricultural Commodity Prices: Empirical Evidence from India”. In: *Journal of Big Data* 12 (Apr. 11, 2025). DOI: 10.1186/s40537-025-01131-8.
- [12] Yongmei Fang et al. “Optimal Forecast Combination Based on Ensemble Empirical Mode Decomposition for Agricultural Commodity Futures Prices”. In: *Journal of Forecasting* (Feb. 17, 2020). DOI: 10.1002/for.2665.
- [13] S. Ray et al. “An ARIMA-LSTM Model for Predicting Volatile Agricultural Price Series with Random Forest Technique”. In: *Appl. Soft Comput.* 149 (Oct. 1, 2023), p. 110939. DOI: 10.1016/j.asoc.2023.110939.
- [14] Murat Sari et al. “Various Optimized Machine Learning Techniques to Predict Agricultural Commodity Prices”. In: *Neural Computing and Applications* 36 (Apr. 16, 2024), pp. 11439–11459. DOI: 10.1007/s00521-024-09679-x.
- [15] Abhibasu Sen and Karabi Dutta Choudhury. “Forecasting the Crude Oil Prices for Last Four Decades Using Deep Learning Approach”. In: *Resources Policy* (Jan. 1, 2024). DOI: 10.1016/j.resourpol.2023.104438.
- [16] Xuedong Liang et al. “Crude Oil Price Prediction Using Deep Reinforcement Learning”. In: *Resources Policy* (Mar. 1, 2023). DOI: 10.1016/j.resourpol.2023.103363.
- [17] A. Varshini, Parthajit Kayal, and Moinak Maiti. “How Good Are Different Machine and Deep Learning Models in Forecasting the Future Price of Metals? Full Sample versus Sub-Sample”. In: *Resources Policy* (May 1, 2024). DOI: 10.1016/j.resourpol.2024.105040.
- [18] Bingzi Jin and Xiaojie Xu. “Machine Learning Gold Price Forecasting”. In: *International Journal of Management Science and Engineering Management* 20 (Jan. 21, 2025), pp. 273–285. DOI: 10.1080/17509653.2025.2453902.

- [19] Sinno Jialin Pan and Qiang Yang. “A Survey on Transfer Learning”. In: *IEEE Transactions on Knowledge and Data Engineering* 22.10 (Oct. 2010), pp. 1345–1359. ISSN: 1558-2191. DOI: 10.1109/TKDE.2009.191.
- [20] Purna Chander Mashetty et al. “Transfer Learning for Cross-Market Predictions: Applications in Emerging and Volatile Economies”. In: *2025 6th International Conference on Data Intelligence and Cognitive Informatics (ICDICI)*. 2025 6th International Conference on Data Intelligence and Cognitive Informatics (ICDICI). July 2025, pp. 621–626. DOI: 10.1109/ICDICI66477.2025.11135015.
- [21] Mingsheng Long et al. “Learning Transferable Features with Deep Adaptation Networks”. In: *Proceedings of the 32nd International Conference on Machine Learning*. International Conference on Machine Learning. PMLR, June 1, 2015, pp. 97–105.
- [22] Yaroslav Ganin et al. *Domain-Adversarial Training of Neural Networks*. May 26, 2016. DOI: 10.48550/arXiv.1505.07818. arXiv: 1505.07818 [stat.ML]. URL: <http://arxiv.org/abs/1505.07818> (visited on 08/11/2026). Pre-published.
- [23] Francis X. Diebold and Roberto S. Mariano. “Comparing Predictive Accuracy”. In: *Journal of Business & Economic Statistics* 13.3 (1995), pp. 253–263. ISSN: 0735-0015. DOI: 10.2307/1392185. JSTOR: 1392185.
- [24] Nicaise Arnaud Kouakan Adanin and D. Balungu. “Predicting Agricultural Commodity Stock Prices: A Comparison of Statistical Methods and Machine Learning Algorithms”. In: *2025 IEEE Ural-Siberian Conference on Biomedical Engineering, Radioelectronics and Information Technology (USBREIT)*. May 12, 2025, pp. 304–307. DOI: 10.1109/usbereit65494.2025.11054062.
- [25] V. Aggarwal et al. “To Determine the Futures Pricing of Metal Commodities Using Deep Learning”. In: *Evergreen* (June 1, 2023). DOI: 10.5109/6793658.
- [26] “Agricultural Market Price Prediction Using Machine Learning and ARIMA Time Series Models: A Review”. In: *Iconic Research and Engineering Journals* (Nov. 14, 2025). DOI: 10.64388/irev9i5-1711906.
- [27] C. Aher et al. “Market Intelligence System for Agricultural Commodity Price Forecasting Using Deep Learning”. In: *2025 6th International Conference on Inventive Research in Computing Applications (ICIRCA)*. June 25, 2025, pp. 1933–1937. DOI: 10.1109/icirca65293.2025.11089639.

- [28] Hind Aldabagh, Xianrong Zheng, and Ravi Mukkamala. “A Hybrid Deep Learning Approach for Crude Oil Price Prediction”. In: *Journal of Risk and Financial Management* (Dec. 6, 2023). DOI: 10.3390/jrfm16120503.
- [29] Ron Alquist, Saroj Bhattarai, and Olivier Coibion. “Commodity-Price Comovement and Global Economic Activity”. In: *Journal of Monetary Economics* 112 (June 1, 2020), pp. 41–56. ISSN: 0304-3932. DOI: 10.1016/j.jmoneco.2019.02.004.
- [30] M. Alshater et al. “Early Warning System to Predict Energy Prices: The Role of Artificial Intelligence and Machine Learning”. In: *Annals of Operations Research* (Aug. 26, 2022), pp. 1–37. DOI: 10.1007/s10479-022-04908-9.
- [31] Hachmi Ben Ameer et al. “Forecasting Commodity Prices: Empirical Evidence Using Deep Learning Tools”. In: *Annals of Operations Research* (Jan. 20, 2023), pp. 1–19. DOI: 10.1007/s10479-022-05076-6.
- [32] Saulius Masteika, Aleksandras Rutkauskas, and Andrea Janes. “Continuous Futures Data Series for Back Testing and Technical Analysis”. In: (Feb. 2012).
- [33] Benjamin R. Auer. “Have Trend-Following Signals in Commodity Futures Markets Become Less Reliable in Recent Years?” In: *Financial Markets and Portfolio Management* 35.4 (Dec. 1, 2021), pp. 533–553. ISSN: 2373-8529. DOI: 10.1007/s11408-021-00385-5.
- [34] G. Avinash et al. “Hidden Markov Guided Deep Learning Models for Forecasting Highly Volatile Agricultural Commodity Prices”. In: *Appl. Soft Comput.* 158 (Apr. 1, 2024), p. 111557. DOI: 10.1016/j.asoc.2024.111557.
- [35] Haithem Awijen et al. “Forecasting Oil Price in Times of Crisis: A New Evidence from Machine Learning versus Deep Learning Models”. In: *Annals of Operations Research* 345 (July 12, 2023), pp. 979–1002. DOI: 10.1007/s10479-023-05400-8.
- [36] Silvio Barra et al. “Deep Learning and Time Series-to-Image Encoding for Financial Forecasting”. In: *IEEE/CAA Journal of Automatica Sinica* 7 (Apr. 28, 2020), pp. 683–692. DOI: 10.1109/jas.2020.1003132.
- [37] S. Bora and Ani Katchova. “Multi-Step Commodity Forecasts Using Deep Learning”. In: *Agricultural Finance Review* (Sept. 2, 2024). DOI: 10.1108/afr-08-2023-0105.
- [38] P. Brignoli et al. “Machine Learning to Predict Grains Futures Prices”. In: *Agricultural Economics* (Mar. 25, 2024). DOI: 10.1111/agec.12828.

- [39] G. Busari and D. Lim. “Crude Oil Price Prediction: A Comparison between AdaBoost-LSTM and AdaBoost-GRU for Improving Forecasting Performance”. In: *Comput. Chem. Eng.* 155 (Dec. 1, 2021), p. 107513. DOI: 10.1016/j.compchemeng.2021.107513.
- [40] Joseph P Byrne, Ryuta Sakemoto, and Bing Xu. “Commodity Price Co-Movement, Heterogeneity and the Time-Varying Impact of Fundamentals”. In: *European Review of Agricultural Economics* 47.2 (Mar. 20, 2020), pp. 499–528. ISSN: 0165-1587. DOI: 10.1093/erae/jbz017.
- [41] Xuesen Cai. “An Attention-Enhanced TimesNet Time Series Model for Predicting the Commodity Price”. In: *Alexandria Engineering Journal* (Oct. 1, 2025). DOI: 10.1016/j.aej.2025.09.018.
- [42] Roy Cerqueti, Valerio Ficcadenti, and Raffaele Mattera. “Investors’ Attention and Network Spillover for Commodity Market Forecasting”. In: *Socio-Economic Planning Sciences* 95 (Oct. 1, 2024), p. 102023. ISSN: 0038-0121. DOI: 10.1016/j.seps.2024.102023.
- [43] Isara Chaowuttisuk and Papis Wongchaisuwat. “Forecasting Thai Agricultural Price: A Deep Learning Approach for Key Commodities”. In: *International Journal of Machine Learning* (Jan. 1, 2026). DOI: 10.18178/ijml.2026.16.1.1185.
- [44] Naina Charak, Dinesh Kumar, and Jasvinder Singh. “Gold Price Forecasting Using ML and DL Models”. In: 2025 IEEE 6th India Council International Subsections Conference (INDISCON). Aug. 21, 2025, pp. 1–6. DOI: 10.1109/indiscon66021.2025.11251948.
- [45] Zhiyuan Chen et al. “Automated Agriculture Commodity Price Prediction System with Machine Learning Techniques”. In: *ArXiv abs/2106.12747* (June 24, 2021). DOI: 10.25046/aj060442.
- [46] Haoyang Chen. “Research on Supply Chain Demand Forecasting Based on Deep Learning”. In: 2025 IEEE 3rd International Conference on Electrical, Automation and Computer Engineering (ICEACE). Dec. 26, 2025, pp. 372–377. DOI: 10.1109/iceace67491.2025.11439886.
- [47] P.B. Deepa and Josephine Daisy. “Deep Learning Based Prediction of Commodity Prices Using LSTM”. In: 2023 4th International Conference on Smart Electronics

- and Communication (ICOSEC). Sept. 20, 2023, pp. 1708–1715. DOI: 10.1109/icosec58147.2023.10276011.
- [48] S. Dineshkumar and Andi Mercy. “Forecasting Agricultural Commodity Prices Using LSTM–GRU Ensemble Networks: Evidence from Indian Crop Markets”. In: 2026 9th International Conference on Inventive Computation Technologies (ICICT). Apr. 15, 2026, pp. 788–793. DOI: 10.1109/icict68280.2026.11511097.
- [49] L. Doroshenko, L. Mastroeni, and A. Mazzocchi. “Wavelet and Deep Learning Framework for Predicting Commodity Prices Under Economic and Financial Uncertainty”. In: *Mathematics* (Apr. 20, 2025). DOI: 10.3390/math13081346.
- [50] K. Douaioui et al. “Machine Learning and Deep Learning Models for Demand Forecasting in Supply Chain Management: A Critical Review”. In: *Applied System Innovation* (Sept. 26, 2024). DOI: 10.3390/asi7050093.
- [51] Yesser Drira, S. Boutouria, and Mouna Boujelbène. “Impact of Key Risk, Green and Behavioral Factors on Agricultural Commodity Price Forecasting amid Volatile Markets: A Deep Learning Approach”. In: *EuroMed Journal of Business* (Sept. 8, 2025). DOI: 10.1108/emjb-01-2025-0020.
- [52] A. Elberawi and Mohamed Belal. “A DEEP LEARNING APPROACH FOR FORECASTING GLOBAL COMMODITIES PRICES”. In: *Future Computing and Informatics Journal* (July 11, 2021). DOI: 10.54623/fue.fcij.6.1.4.
- [53] El Filali et al. “Machine Learning Applications in Supply Chain Management: A Deep Learning Model Using an Optimized LSTM Network for Demand Forecasting”. In: *International Journal of Intelligent Engineering and Systems* (Apr. 30, 2022). DOI: 10.22266/ijies2022.0430.42.
- [54] Hamed Ghoddusi, Germán Creamer, and N. Rafizadeh. “Machine Learning in Energy Economics and Finance: A Review”. In: *Energy Economics* (June 1, 2019). DOI: 10.1016/j.eneco.2019.05.006.
- [55] Marta Giampietro, Massimo Guidolin, and Manuela Pedio. “Estimating Stochastic Discount Factor Models with Hidden Regimes: Applications to Commodity Pricing”. In: *European Journal of Operational Research* 265.2 (Mar. 1, 2018), pp. 685–702. ISSN: 0377-2217. DOI: 10.1016/j.ejor.2017.07.045.

- [56] Sofia Giantsidi and Claudia Tarantola. “Deep Learning for Financial Forecasting: A Review of Recent Trends”. In: *International Review of Economics & Finance* (Nov. 1, 2025). DOI: 10.1016/j.iref.2025.104719.
- [57] Harsh Gupta et al. “Forecasting Commodity Prices Using Machine Learning”. In: *International Journal of Scientific Research in Science and Technology* (Jan. 1, 2024). DOI: 10.32628/ijrst52411110.
- [58] Asmaa Hamad et al. “Machine Learning Techniques for Accurate Forecasting of Commodity Prices”. In: 2025 3rd International Conference on Business Analytics for Technology and Security (ICBATS). May 1, 2025, pp. 1–7. DOI: 10.1109/icbats66542.2025.11258485.
- [59] F. Hamouda, Nadia Arfaoui, and Md Naeem. “Forecasting Energy Commodity Prices Amidst Worldwide Energy Transitions Using Artificial Intelligence Models”. In: *The Energy Journal* 46 (June 27, 2025), pp. 215–244. DOI: 10.1177/01956574251340012.
- [60] A. Hannon, Ch. Paramaiah, and Amani Bin Tareef. “Forecasting Commodities Prices Using Machine Learning”. In: 2025 International Conference on Business Intelligence for Technology Innovation (ICBITI). Apr. 16, 2025, pp. 1–8. DOI: 10.1109/icbiti65527.2025.11500869.
- [61] Wei Hao and Linh Pham. “Dynamic Connectedness in the Higher Moments between Clean Energy and Oil Prices”. In: *Energy Economics* 140 (Dec. 1, 2024), p. 107987. ISSN: 0140-9883. DOI: 10.1016/j.eneco.2024.107987.
- [62] N. Harshith, Prity Kumari, and R. Pundir. “Memory Based Neural Network for Cumin Price Forecasting in Gujarat, India”. In: *Journal of Agriculture and Food Research* (Jan. 1, 2024). DOI: 10.1016/j.jafr.2024.101020.
- [63] Gabriel Paes Herrera et al. “Long-Term Forecast of Energy Commodities Price Using Machine Learning”. In: *Energy* (July 15, 2019). DOI: 10.1016/j.energy.2019.04.077.
- [64] Hasnain Iftikhar et al. “Daily Crude Oil Prices Forecasting Using a Novel Hybrid Time Series Technique”. In: *IEEE Access* 13 (2025), pp. 98822–98836. DOI: 10.1109/access.2025.3574788.
- [65] Bingzi Jin and Xiaojie Xu. “China Commodity Price Index (CCPI) Forecasting via the Neural Network”. In: *International Journal of Financial Engineering* (Feb. 21, 2025). DOI: 10.1142/s2424786325500033.

- [66] Bingzi Jin and Xiaojie Xu. “Machine Learning Natural Gas Price Predictions”. In: *Energy Science* (Jan. 10, 2025). DOI: 10.1142/s2972379525500012.
- [67] Bingzi Jin, Xiaojie Xu, and Yun Zhang. “Peanut Oil Price Change Forecasts through the Neural Network”. In: *foresight* (Jan. 16, 2025). DOI: 10.1108/fs-01-2023-0016.
- [68] Bingzi Jin and Xiaojie Xu. “Price Forecasting through Neural Networks for Crude Oil, Heating Oil, and Natural Gas”. In: *Measurement: Energy* (Feb. 1, 2024). DOI: 10.1016/j.meaeene.2024.100001.
- [69] Abisheik K. “Agriculture Commodity Price Forecasting”. In: *International Journal for Research in Applied Science and Engineering Technology* (May 31, 2026). DOI: 10.22214/ijraset.2026.82392.
- [70] Ilias Kampouris, Charilaos Mertzanis, and Aristeidis Samitas. “Natural Disaster Shocks and Commodity Market Volatility: A Machine Learning Approach”. In: *Pacific-Basin Finance Journal* (Dec. 1, 2024). DOI: 10.1016/j.pacfin.2024.102618.
- [71] N. Karmarkar. “Time Series Analysis for Commodity Price Forecasting”. In: *International Journal of Commerce, Finance and Digital Economy* (Jan. 1, 2024). DOI: 10.67228/3071642x/ijcfde-2024pii7k1d.
- [72] Baljinder Kaur. “Predictive Analytics Models for Commodity Market: A Literature Review”. In: *International Journal of Intelligent Systems and Applications in Engineering* 12.3 (Aug. 30, 2024), pp. 4525–4535. ISSN: 2147-6799.
- [73] Ahhyun Kim, Doojin Ryu, and Alexander Webb. “Forecasting Oil Futures Markets Using Machine Learning and Seasonal Trend Decomposition”. In: *Investment Analysts Journal* 54 (Oct. 24, 2024), pp. 205–218. DOI: 10.1080/10293523.2024.2405294.
- [74] Prity Kumari et al. “Recurrent Neural Network Architecture for Forecasting Banana Prices in Gujarat, India”. In: *PLOS ONE* 18 (June 15, 2023). DOI: 10.1371/journal.pone.0275702.
- [75] Ravindra Kumar et al. “Precision Price Forecasting for Agricultural Commodities Using AI-Driven Temporal Models”. In: 2025 2nd Asia Pacific Conference on Innovation in Technology (APCIT). Sept. 19, 2025, pp. 1–7. DOI: 10.1109/apcit65661.2025.11411508.

- [76] M. L. and A. Mishra. “Forecasting Spot Prices of Agricultural Commodities in India: Application of Deep-Learning Models”. In: *Intell. Syst. Account. Finance Manag.* 28 (Jan. 28, 2021), pp. 72–83. DOI: 10.1002/isaf.1487.
- [77] S. Lahmiri. “Fossil Energy Market Price Prediction by Using Machine Learning with Optimal Hyper-Parameters: A Comparative Study”. In: *Resources Policy* (May 1, 2024). DOI: 10.1016/j.resourpol.2024.105008.
- [78] Ngoc-Bao-Van Le, Yeong-Seok Seo, and Jun-Ho Huh. “AgTech: Volatility Prediction for Agricultural Commodity Exchange Trading Applied Deep Learning”. In: *IEEE Access* 12 (2024), pp. 153898–153910. DOI: 10.1109/access.2024.3479868.
- [79] Dwi Lestari et al. “Machine Learning-Based Forecasting of Agricultural Commodity Prices Using Ensemble Models”. In: *Journal of Human, Earth, and Future* (Dec. 1, 2025). DOI: 10.28991/hef-2025-06-04-09.
- [80] Pan Li, Jingyi Liu, and Lei Qiu. “Deep Learning Methods for Demand Forecasting and Inventory Optimization in Modern Supply Chains”. In: *Asian Business Research Journal* (Mar. 10, 2026). DOI: 10.55220/2576-6759.v11i3.906.
- [81] Xuerong Li, Wei Shang, and Shouyang Wang. “Text-Based Crude Oil Price Forecasting: A Deep Learning Approach”. In: *International Journal of Forecasting* (Oct. 1, 2019). DOI: 10.1016/j.ijforecast.2018.07.006.
- [82] Alexandre Lucas et al. “Price Forecasting for the Balancing Energy Market Using Machine-Learning Regression”. In: *Energies* (Oct. 16, 2020). DOI: 10.3390/en13205420.
- [83] A. Mahto et al. “Short-Term Forecasting of Agriculture Commodities in Context of Indian Market for Sustainable Agriculture by Using the Artificial Neural Network”. In: *Journal of Food Quality* 2021 (May 4, 2021), pp. 1–13. DOI: 10.1155/2021/9939906.
- [84] N. Manikandan et al. “Forecasting Agricultural Commodity Prices Using Machine Learning”. In: 2025 6th International Conference on Data Intelligence and Cognitive Informatics (ICDICI). July 9, 2025, pp. 1827–1832. DOI: 10.1109/icdici66477.2025.11135297.
- [85] R. Manogna, Vijay Dharmaji, and S. Sarang. “Enhancing Agricultural Commodity Price Forecasting with Deep Learning”. In: *Scientific Reports* 15 (July 1, 2025). DOI: 10.1038/s41598-025-05103-z.

- [86] Simone Merello et al. “Ensemble Application of Transfer Learning and Sample Weighting for Stock Market Prediction”. In: *2019 International Joint Conference on Neural Networks (IJCNN)*. 2019 International Joint Conference on Neural Networks (IJCNN). July 2019, pp. 1–8. DOI: 10.1109/IJCNN.2019.8851938.
- [87] Mostafa Moghadam. *Forecasting Asset Returns with Sentiment-Enhanced FI-GARCH Models*. Dec. 23, 2025. DOI: 10.2139/ssrn.5943491. Social Science Research Network: 5943491. URL: <https://papers.ssrn.com/abstract=5943491> (visited on 06/13/2026). Pre-published.
- [88] Manas Kumar Mohanty, P. Thakurta, and S. Kar. “Agricultural Commodity Price Prediction Model: A Machine Learning Framework”. In: *Neural Computing and Applications* 35 (Apr. 6, 2023), pp. 15109–15128. DOI: 10.1007/s00521-023-08528-7.
- [89] Tashreef Muhammad et al. “A Benchmark of Classical and Deep Learning Models for Agricultural Commodity Price Forecasting on A Novel Bangladeshi Market Price Dataset”. In: *ArXiv abs/2604.06227* (Mar. 27, 2026). DOI: 10.48550/arxiv.2604.06227.
- [90] Tushar Babulal Nandeshwar, Akanksha Sharma, and C. Verma. “Assessing Deep Learning Techniques for Commodity Price Prediction: A Comparative Assessment Across Energy, Precious Metal, and Agricultural Futures”. In: *2025 International Conference on Computational, Communication and Information Technology (ICCCIT)*. Feb. 7, 2025, pp. 74–79. DOI: 10.1109/icccit62592.2025.10927995.
- [91] G. Nayak et al. “Exogenous Variable Driven Deep Learning Models for Improved Price Forecasting of TOP Crops in India”. In: *Scientific Reports* 14 (July 26, 2024), DOI: 10.1038/s41598-024-68040-3.
- [92] G. Nayak et al. “N-BEATS Deep Learning Architecture for Agricultural Commodity Price Forecasting”. In: *Potato Research* 68 (Sept. 12, 2024), pp. 1437–1457. DOI: 10.1007/s11540-024-09789-y.
- [93] Seyed Mahdi Mirhoseini Nejad et al. “Multispectral Crop Yield Prediction Using 3D-Convolutional Neural Networks and Attention Convolutional LSTM Approaches”. In: *IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing* 16 (2023), pp. 254–266. DOI: 10.1109/jstars.2022.3223423.

- [94] R. Paul et al. "Can Deep Learning Models Enhance the Accuracy of Agricultural Price Forecasting? Insights From India". In: *Intell. Syst. Account. Finance Manag.* 32 (Mar. 1, 2025). DOI: 10.1002/isaf.70002.
- [95] Sruthi Paul et al. "Deep Learning-Based Price Forecasting in Energy Markets: A Systematic Literature Review and Future Research Directions". In: *Engineering Research Express* 8 (June 2, 2026). DOI: 10.1088/2631-8695/ae76d8.
- [96] R. Paul et al. "Machine Learning Techniques for Forecasting Agricultural Prices: A Case of Brinjal in Odisha, India". In: *PLoS ONE* 17 (July 6, 2022). DOI: 10.1371/journal.pone.0270553.
- [97] Hn Raghu and Aluregowda. "Unlocking Future Commodities Markets: Innovative Approaches to Price Prediction Using Supervised Machine Learning". In: *Shanlax International Journal of Arts, Science and Humanities* (Oct. 28, 2024). DOI: 10.34293/sijash.v12is1-oct.8295.
- [98] B. Rajana et al. "Hybrid Econometric and Deep Learning Framework for Forecasting Global Commodity Price Dynamics: Evidence from Multi-Commodity Time Series (2000-2026)". In: *Global Academic Journal of Economics and Business* (Apr. 20, 2026). DOI: 10.36348/gajeb.2026.v08i02.010.
- [99] Soumik Ray et al. "A Random Forest-Convolutional Neural Network Deep Learning Model for Predicting the Wholesale Price Index of Potato in India". In: *Potato Research* 68 (May 24, 2024), pp. 263–279. DOI: 10.1007/s11540-024-09736-x.
- [100] V. R, Yashaswini G, and Kavyashree J. "Predictive Analytics for Agricultural Commodity Markets Using Machine Learning". In: *INTERNATIONAL JOURNAL OF SCIENTIFIC RESEARCH IN ENGINEERING AND MANAGEMENT* (Aug. 18, 2025). DOI: 10.55041/ijsrem51858.
- [101] Natasha Saeed et al. "Intelligent Decision Making for Commodities Price Prediction: Opportunities, Challenges and Future Avenues". In: *Computational Economics* 66 (Jan. 7, 2025), pp. 3781–3839. DOI: 10.1007/s10614-024-10837-5.
- [102] M. Sandeep et al. "AI-ML Based Price Prediction Model for Agri-Horticultural Commodities". In: *International Journal of Research Publication and Reviews* (Mar. 1, 2025). DOI: 10.55248/gengpi.6.0325.1179.

- [103] O. Sezer, M. Gudelek, and A. Ozbayoglu. “Financial Time Series Forecasting with Deep Learning : A Systematic Literature Review: 2005-2019”. In: *ArXiv abs/1911.13288* (Nov. 29, 2019). DOI: 10.1016/j.asoc.2020.106181.
- [104] F. Shaheen et al. “Deep Learning-Enabled Cherry Price Forecasting and Real-Time System Deployment across Multi-Market Supply Chains in India”. In: *Scientific Reports* 16 (Dec. 10, 2025). DOI: 10.1038/s41598-025-30980-9.
- [105] P. Shobha et al. “Intelligent Forecasting of Agricultural Markets Using Ensemble of ARIMA, Prophet, LSTM, and GARCH”. In: 2025 6th International Conference on Communication, Computing & Industry 6.0 (C2I6). Dec. 5, 2025, pp. 1–6. DOI: 10.1109/c2i666499.2025.11366981.
- [106] Swapnil Shrivastava, Supriya Pal, and Ranjeet Walia. “Market Intelligence for Agricultural Commodities Using Forecasting and Deep Learning Techniques”. In: (Dec. 17, 2019), pp. 193–205. DOI: 10.1007/978-3-030-37188-3_12.
- [107] A. Soladoye et al. “Hybridized LSTM-GRU Model for Forecasting the Prices of Crude Oil”. In: *FUOYE Journal of Engineering and Technology* (Aug. 2, 2024). DOI: 10.4314/fuoyejet.v9i2.14.
- [108] Feihu Sun et al. “Agricultural Product Price Forecasting Methods: A Review”. In: *Agriculture* 13.9 (Sept. 2023), p. 1671. ISSN: 2077-0472. DOI: 10.3390/agriculture13091671.
- [109] Mohammad Tami and A. Owda. “Efficient Commodity Price Forecasting Using Long Short-Term Memory Model”. In: *IAES International Journal of Artificial Intelligence (IJ-AI)* (Mar. 1, 2024). DOI: 10.11591/ijai.v13.i1.pp994-1004.
- [110] Mohammed Gadafi Tamimu et al. “Attention-Based Deep Learning Hybrid Model for Cash Crop Price Forecasting: Evidence from Global Futures Markets with Implications for West Africa”. In: *Applied Sciences* (Feb. 5, 2026). DOI: 10.3390/app16031600.
- [111] Yogesh Thakre et al. “Commodity Price Prediction Using Machine Learning: Insights, Challenges, and Prospect”. In: 2025 International Conference on Recent Advances in Electrical, Electronics, Ubiquitous Communication, and Computational Intelligence (RAEEUCCI). Apr. 23, 2025, pp. 1–7. DOI: 10.1109/raeeucci63961.2025.11048306.

- [112] Asterios Theofilou et al. “Predicting Prices of Staple Crops Using Machine Learning: A Systematic Review of Studies on Wheat, Corn, and Rice”. In: *Sustainability* (June 13, 2025). DOI: 10.3390/su17125456.
- [113] Kais Tissaoui and Nadia Azibi. “Sentiment Momentum and Regime-Dependent Predictability in Crude Oil Markets: A Markov Switching Analysis with Higher-Order VIX and MOVE Derivatives”. In: *International Review of Economics and Finance* 108 (2026). ISSN: 1059-0560. DOI: 10.1016/j.iref.2026.105197.
- [114] G. T. “A Review of Deep Learning Models for Price Prediction in Agricultural Commodities”. In: *Economic Affairs* (Mar. 25, 2023). DOI: 10.46852/0424-2513.1.2023.25.
- [115] B. Varshini et al. “CNN-transformer Hybrid Deep Learning Model for Agricultural Commodities Price Forecasting”. In: *Discover Computing* (Mar. 8, 2026). DOI: 10.1007/s10791-026-09966-2.
- [116] Zilin Wang et al. “Climate and Environmental Data Contribute to the Prediction of Grain Commodity Prices Using Deep Learning”. In: *Journal of Sustainable Agriculture and Environment* (June 6, 2023). DOI: 10.1002/sae2.12041.
- [117] Le Wang and Boyuan Zhang. “The Promise of Time-Series Foundation Models for Agricultural Forecasting: Evidence from Commodity Prices”. In: (Jan. 10, 2026). DOI: 10.48550/arxiv.2601.06371.
- [118] Hlaing Hlaing Win, Kyawt Kyawt San, and Khin Ei Ei Chaw. “Deep Learning for Commodity Price Forecasting: LSTM, BiLSTM, GRU Comparison with Variable Window Sizes and Activation Functions”. In: 2024 5th International Conference on Advanced Information Technologies (ICAIT). Nov. 7, 2024, pp. 1–6. DOI: 10.1109/icaait65209.2024.10754939.
- [119] Hlaing Hlaing Win and T. Zin. “Time Series Forecasting of Commodity Prices Using ARIMA, Holt-Winters and LSTM”. In: 2025 6th International Conference on Advanced Information Technologies (ICAIT). Nov. 3, 2025, pp. 1–6. DOI: 10.1109/icaait68809.2025.11236779.
- [120] Xiaojie Xu and Yun Zhang. “Commodity Price Forecasting via Neural Networks for Coffee, Corn, Cotton, Oats, Soybeans, Soybean Oil, Sugar, and Wheat”. In: *Intell. Syst. Account. Finance Manag.* 29 (Aug. 3, 2022), pp. 169–181. DOI: 10.1002/isaf.1519.

- [121] Zili Xu et al. “Using Econometric and Machine Learning Models to Forecast Crude Oil Prices: Insights from Economic History”. In: *Resources Policy* (June 1, 2023). DOI: 10.1016/j.resourpol.2023.103614.
- [122] Hao Yang, Jie Yang, and Yun Feng. “Climate Physical Risks and the Vulnerability of Global Agricultural Commodities”. In: *Economics Letters* 258 (Jan. 1, 2026), p. 112748. ISSN: 0165-1765. DOI: 10.1016/j.econlet.2025.112748.
- [123] Kailing Yang et al. “Predicting Energy Prices Based on a Novel Hybrid Machine Learning: Comprehensive Study of Multi-Step Price Forecasting”. In: *Energy* (Apr. 1, 2024). DOI: 10.1016/j.energy.2024.131321.
- [124] Dabin Zhang et al. “Forecasting Agricultural Commodity Prices Using Model Selection Framework With Time Series Features and Forecast Horizons”. In: *IEEE Access* 8 (Feb. 4, 2020), pp. 28197–28209. DOI: 10.1109/access.2020.2971591.
- [125] Yang Zhao, Jianping Li, and Lean Yu. “A Deep Learning Ensemble Approach for Crude Oil Price Forecasting”. In: *Energy Economics* 66 (Aug. 1, 2017), pp. 9–16. DOI: 10.1016/j.eneco.2017.05.023.

*not cited yet

List of Acronyms

- Adam** Adaptive Moment Estimation
- ARIMA** Autoregressive Integrated Moving Average
- BiLSTM** Bidirectional Long Short-Term Memory
- CNN** Convolutional Neural Network
- DANN** Domain-Adversarial Neural Network
- ESN** Echo State Model
- GARCH** Generalized Autoregressive Conditional Heteroskedasticity
- GRU** Gated Recurrent Unit
- LSTM** Long Short-Term Memory
- MAE** Mean Absolute Error
- MLP** Multilayer Perceptron
- MMD** Maximum Mean Discrepancy
- NN** Neural Network
- RBF** Radial Basis Function
- RKHS** Reproducing Kernel Hilbert Space
- RMSE** Root Mean Squared Error
- RNN** Recurrent Neural Network
- SVM** Support Vector Machine
- SVR** Support Vector Regression
- TCN** Temporal Convolutional Network
- XGBoost** Extreme Gradient Boosting

Appendix

Full Results Tables

Model Architecture Specifications

Data Sources and Preprocessing

Statistical Tests